# Re-estimation-verified selection of small parameter subsets for systems biology models

Mónica Miranda R.[1*], Carlos Salazar[1], César Martín[1], Adriana Aguirre[1]

[1*]Faculty of Engineering, Control and Automation Systems Group, Escuela Superior Politécnica del Litoral (ESPOL), Guayaquil, Ecuador.

*Corresponding author(s). E-mail(s): mmmirand@espol.edu.ec;
Contributing authors: asalazar@espol.edu.ec; cmartin@espol.edu.ec; aaaguirr@espol.edu.ec;

**Abstract**

Systems biology models have more rate constants than data can determine. Starting from published values, our method ranks biological parameters by output sensitivity, rejects redundant ones and eliminates parameters backwards to an imposed minimum of two, accepting a subset only if its re-estimation, together with the measurement scaling and offset factors, reproduces, within a set tolerance, the output response to random parameter perturbations. In 31 of 34 benchmark problems with real data, a subset of usually two biological parameters passed at ±5% (one sufficed in 15 of 27), confirmed by nonlinear re-estimation, and 22 remained admissible up to ±50% in both robustness designs (five evaluated to first order). The scaling and offset factors alone passed in 9 of 28 problems and random subsets of the same size in 16, so the test discriminates only in some problems. With the others at published values, all selected parameters were practically identifiable in 25 problems. Selected on training data, our subsets predicted synthetic data near the published values more often than pivoted QR and profile likelihood, and not detectably differently from Brun's collinearity index. The general theorems behind the method are verified in Lean 4; their application to each result is checked numerically.

**Keywords:** parameter reduction, practical identifiability, sensitivity analysis, ordinary differential equations, systems biology, formal verification

# Introduction

Ordinary differential equation (ODE) models are central to quantitative biology, pharmacology and epidemiology, where they describe how interacting molecular species or populations change over time as a function of kinetic parameters $\theta$[1]. Using a model for prediction requires estimating $\theta$ from measurements, and this becomes difficult as models grow. The main obstacle is non-identifiability: different parameter values may produce the same measured output, because of model symmetries (structural non-identifiability) or because the data are insufficient (practical non-identifiability)[2–6]. The likelihood then has flat ridges on which optimisers stall and confidence intervals become unreliable[7–9]. Models in public benchmark collections routinely contain 20–300 parameters[10].

The sloppy-model phenomenon offers a structural explanation[11, 12]: the outputs are typically sensitive to a few "stiff" parameter directions and almost insensitive to many "sloppy" ones[13], which suggests that a small subset of parameters may suffice to describe what is actually measured. Global sensitivity indices[14, 15] and profile likelihoods[7, 16, 17] rank parameters by importance or identifiability, and subset-selection methods based on the sensitivity or Fisher information matrix additionally avoid redundant parameters[18]: Fisher-information criteria[19, 20], collinearity indices[21, 22], orthogonalization[23–25], eigenvalue methods[26] and pivoted QR[28, 29], several of which have been compared on signalling and physiological models[26, 30]. All of them return a subset, but, to our knowledge, none tests whether re-estimating the selected subset reproduces the outputs of the full model when all parameters change, and the validity of statements such as "the discarded parameters can be fixed" is rarely made explicit.

This work was motivated by a 22-parameter model of Bandura's social cognitive theory[31–34], in which many stable parameter sets fitted the same data with widely different quality[35].

Here we present a reduction method that works directly on the measured outputs defined in the PEtab format[36], with their experimental conditions, transformations and noise levels. It separates biological parameters from calibration parameters (scalings and offsets of the measurement), which are always re-estimated, and selects only among the biological ones: candidates are ranked by sensitivity, redundant ones are rejected by conditioning tests, the reduced model is accepted only if re-estimating it reproduces the full model under random parameter perturbations, and backward elimination removes every parameter whose removal keeps the subset admissible under the rule $|S| \geq 2$. We evaluate the method on 34 benchmark problems with real data[10, 36], on held-out experiments and synthetic data, against four existing criteria and null baselines under a common protocol, and we report its robustness and computational cost. Our contributions are threefold: (i) a selection criterion that accepts a subset only after verifying by re-estimation that it reproduces the full model on the measured outputs; (ii) a benchmark-wide comparison with four criteria and two null baselines under a common protocol, which shows both where the test separates subsets and where it does not; and (iii) general theorems behind each step verified in the Lean 4 proof assistant with Mathlib[37, 38], with an explicit statement of what is proved in general and what is only checked numerically for each result.

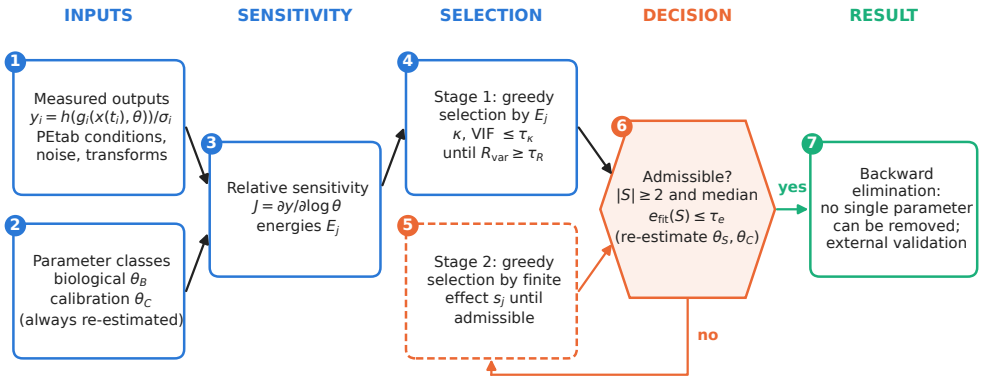



**Fig. 1 Overview of the reduction method.** Biological parameters are ranked by their effect on the measured outputs and selected with conditioning control ($\kappa$ and VIF thresholds; stage 1, or stage 2 if needed, dashed). A subset is accepted when re-estimating it reproduces the full model; backward elimination then removes every parameter that is individually dispensable (the result is inclusion-minimal if the criterion is monotone; Theorem 2), and the reduced model is validated on held-out data.

# Results

## Overview of the method

Figure 1 summarises the procedure; symbols are defined in Methods (*Measured outputs and parameter classes* to *Admissibility and re-estimation*) and listed in Supplementary Table S2. The input is a PEtab problem: an SBML model, the measured observables with their conditions and noise standard deviations, and the published parameter values $\theta_0$. Outputs are scaled by the noise; only biological parameters $\theta_B$ (those of the SBML model) are candidates for selection, and calibration parameters $\theta_C$ (scalings and offsets) are always re-estimated with the selected ones.

The relative sensitivity matrix $J = \partial y/\partial \log \theta$ is computed once at $\theta_0$. Stage 1 adds biological parameters in decreasing order of their sensitivity energy $E_j = \|J_{:,j}\|^2$, accepting a candidate only if the condition number $\kappa$ and the largest variance inflation factor (VIF) stay below a threshold, until a fraction $\tau_R$ of the energy is captured. A subset $S$ is admissible if, in 15 random scenarios that perturb all parameters by $\pm 5\%$, re-estimating $\theta_S$ and $\theta_C$ with the other parameters at $\theta_0$ gives a median relative fitting error $e_{\text{fit}} \leq \tau_e$, and $|S| \geq 2$; $\tau_e = \sqrt{1 - c^2} = 0.436$ ($c = 0.9$) is the error left by rescaling an output change whose cosine similarity with the full change is $c$ (Theorem 1). If stage 1 fails, stage 2 orders the parameters by their finite-perturbation effect and stops at the first admissible subset; backward elimination then removes any parameter whose removal keeps the subset admissible (Theorem 2).

## Benchmark: a few biological parameters suffice in most models

We applied the method to 34 PEtab benchmark problems with real measurements[10, 36] from signalling, metabolism, gene regulation, cell cycle,

immunology and epidemiology, with 3 to 294 parameters and 16 to 27,132 measurements (Table 1; Methods, *Models, data and simulation*; Supplementary Note *System-by-system summary*). An admissible subset was found in 31 (Fig. 2); in the six with the most expensive simulations admissibility was decided to first order and confirmed by nonlinear re-estimation of the final subset (in Lucarelli with a margin of $4 \times 10^{-5}$; Methods, *Admissibility and re-estimation*, *First-order $e_{\mathrm{fit}}$ and nonlinear check*). In 28 problems the subset has two biological parameters, the minimum imposed by the method ($k_{\min} = 2$), although one already reaches the threshold in 15 of 27 tested problems (Methods, *Settings*, *Single-parameter subsets*); Zhao, Bruno and Lucarelli need three, five and ten. With the calibration parameters, the reduced problem has a median of 78% fewer parameters than the PEtab problem (33–97%; Fig. 2a), or 73% (0–97%) without the fixed noise parameters, the 0 being Perelson, whose two biological parameters are both selected; among biological parameters the reduction is 86%.

All 31 subsets satisfy the conditioning thresholds of the selection on the biological columns (maximum $\kappa = 6.5$, VIF$= 9.8$; with the calibration parameters the VIF exceeds 10 in ten problems; Methods, *Sensitivities, energy and conditioning*). In 17 the median cosine also reaches $c$, so that a single rescaling suffices; in Alkan and Perelson the captured energy reaches $\tau_R$, but the corresponding formal bound requires a model verified in Lean, which these two are not (levels of guarantee in Table 1). Stage 2 was needed in 24 problems (Methods, *Selection*). Of the three failures, Crauste has no local reduction because 1% changes make its solution blow up within the measured horizon (Methods, *Models, data and simulation*, *Crauste*), and in Beer and Lang the search stopped at the conditioning limit with 18 and 13 parameters; in Lang other thresholds give admissible pairs, so the failure is one of the greedy search (Methods, *Settings*).

Re-estimation is essential: with the selected parameters merely set to their perturbed values, the median error $e_{\mathrm{rel}}$ exceeds $\tau_e$ in 23 of the 31 problems (Fig. 2b), whereas after re-estimation the median $e_{\mathrm{fit}}$ is 15% (maximum 43.585%, Lucarelli). The selected parameters compensate for the perturbations of all the others, the property needed when the reduced model is fitted to data.

## The selected parameters are also identifiable

A subset that reproduces the full model need not be determined by the data. Profiling every selected parameter in the reduced problem ($S \cup \Theta_C$ free, the others at $\theta_0$; Methods, *Identifiability of the selected parameters*), all selected parameters were practically identifiable, with a 95% interval bounded on both sides, in 25 of the 31 problems (65 of 74 parameters; Table 1, column Ident.), or in 23 problems (60 parameters) with the threshold scaled by the misfit $\chi^2$/d.f.. These intervals are conditional on the published values of the fixed parameters and therefore understate the uncertainty. The exceptions were Fujita (a rate constant confounded with a scaling factor), Giordano and Perelson (noise far above the residuals), Okuonghae and Oliveira (a profile could not be simulated) and Smith (optimum at the bound). With all parameters free, only 20 of the 52 evaluable selected parameters were identifiable: the data do not determine these models, but once the others are fixed

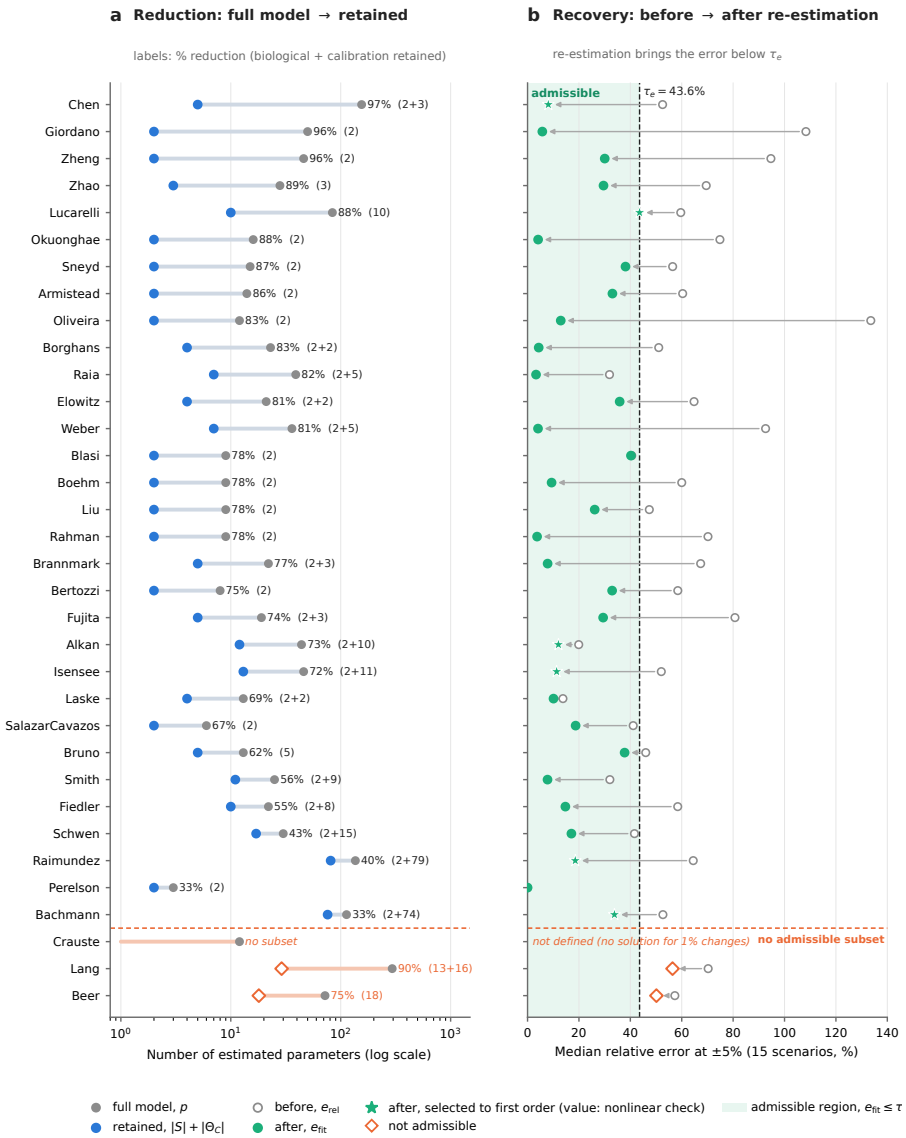



**Fig. 2 Parameter reduction and recovery of the full-model output in 34 benchmark problems.** Problems are ordered by the reduction achieved; the three without an admissible subset (Crauste, Lang and Beer) are at the bottom. **a**, Estimated parameters of the full model ($p$, grey, including the noise parameters, which are kept fixed) and parameters re-estimated in the reduced model (selected biological plus calibration parameters, blue); the percentage is the reduction relative to $p$. Crauste has no admissible subset, so no reduction is shown. **b**, Median relative error over 15 random scenarios at $\pm 5\%$ before ($e_{\mathrm{rel}}$, open circles) and after re-estimating the retained parameters ($e_{\mathrm{fit}}$, green; stars, subsets selected with the first-order $e_{\mathrm{fit}}$, shown with the value of the nonlinear re-estimation of the final subset; for Beer and Lang, first-order value). The shaded region is the admissible region, $e_{\mathrm{fit}} \leq \tau_e = 43.6\%$.

**Table 1 Per-problem results.** $p$, number of parameters estimated in the PEtab problem, including noise parameters, which are kept fixed (Supplementary Table S3); $m$, number of measurements; $|\Theta_C|$, calibration parameters (always re-estimated); $|S|$, selected biological parameters; Stage, stage that produced the subset; $e_{\mathrm{rel}}$ and $e_{\mathrm{fit}}$, median relative error over 15 scenarios at ±5% without and with re-estimation (admissibility: $e_{\mathrm{fit}} \leq 43.6\%$; italic, not admissible; $^*$, subset selected with the first-order $e_{\mathrm{fit}}$; the value shown is the nonlinear re-estimation of the final subset, except for Beer and Lang, which have no admissible subset (first-order value); $^\dagger$, not defined because the model has no solution on the measured horizon for 1% parameter changes); Pattern, behaviour of the median $e_{\mathrm{fit}}$ for perturbation sizes from ±1% to ±50% (A, admissible at all sizes; B, admissible up to at least ±10%, first failing at ±20% or more and not recovering; C, irregular; D, not admissible at ±10% or less; –, not applicable); Level, level of guarantee (1, $e_{\mathrm{fit}} \leq \tau_e$; 2, in addition median $\cos\Delta \geq c$; 3, in addition $R_{\mathrm{var}} \geq \tau_R$); Ident., selected parameters whose profile likelihood in the reduced problem exceeds $\Delta\chi^2 = 3.84$ on both sides within the bounds, out of $|S|$, and in parentheses those for which the hypotheses of the Lean identifiability theorem were satisfied at all evaluated profile points (a numerical check, not a certificate; Table 3); Test $\chi^2/n$, mean squared noise-scaled residual on held-out experiments for the reduced and the full model, both re-estimated on the training experiments (values near 1 mean errors of the size of the noise; –, not evaluable or not run).

| System | $p$ | $m$ | $\lvert\Theta_C\rvert$ | $\lvert S\rvert$ | Stage | $e_{\mathrm{rel}}$ (%) | $e_{\mathrm{fit}}$ (%) | Level | Pattern | Ident. (Lean) | Test $\chi^2/n$ red. / full |
|---|---|---|---|---|---|---|---|---|---|---|---|
| Alkan | 44 | 1733 | 10 | 2 | 1 | 20 | 12.0$^*$ | 3 | A | 2/2 (1) | 2.75 / 3.37 |
| Armistead | 14 | 58 | 0 | 2 | 2 | 60 | 33.0 | 1 | B | 2/2 (2) | 1.22 / 1.37 |
| Bachmann | 113 | 541 | 74 | 2 | 2 | 53 | 33.8$^*$ | 1 | A | 2/2 (0) | 4.77 / 7.35 |
| Beer | 72 | 27132 | 0 | 18 | 2 | 57 | *50.1*$^*$ | – | – | – | – |
| Bertozzi | 8 | 22 | 0 | 2 | 2 | 58 | 32.9 | 2 | B | 2/2 (0) | – |
| Blasi | 9 | 252 | 0 | 2 | 1 | 41 | 40.2 | 2 | B | 2/2 (2) | – |
| Boehm | 9 | 48 | 0 | 2 | 2 | 60 | 9.4 | 2 | A | 2/2 (2) | 2.57 / 11.22 |
| Borghans | 23 | 111 | 2 | 2 | 2 | 51 | 4.3 | 1 | A | 2/2 (0) | 0.74 / 0.74 |
| Brannmark | 22 | 43 | 3 | 2 | 2 | 67 | 7.8 | 1 | A | 2/2 (0) | 5.96 / 6,277 |
| Bruno | 13 | 77 | 0 | 5 | 1 | 46 | 37.8 | 2 | A | 5/5 (5) | 0.91 / 208 |
| Chen | 155 | 120 | 3 | 2 | 2 | 53 | 8.1$^*$ | 1 | A | 2/2 (0) | 20.88 / 17.71 |
| Crauste | 12 | 21 | 0 | – | 2 | – | *n.d.*$^\dagger$ | – | – | – | – |
| Elowitz | 21 | 58 | 2 | 2 | 1 | 65 | 35.8 | 1 | B | 2/2 (0) | 1.50 / 3.89 |
| Fiedler | 22 | 72 | 8 | 2 | 2 | 58 | 14.7 | 1 | A | 2/2 (0) | 11.21 / 9.32 |
| Fujita | 19 | 144 | 3 | 2 | 2 | 81 | 29.4 | 1 | A | 0/2 (0) | 5.94 / 8.03 |
| Giordano | 50 | 313 | 0 | 2 | 2 | 108 | 5.8 | 2 | A | 0/2 (0) | <0.01 / <0.01 |
| Isensee | 46 | 687 | 11 | 2 | 1 | 52 | 11.3$^*$ | 1 | A | 2/2 (0) | 1.08 / 1.08 |
| Lang | 294 | 9600 | 16 | 13 | 2 | 70 | *56.4*$^*$ | – | – | – | – |
| Laske | 13 | 42 | 2 | 2 | 1 | 14 | 10.1 | 2 | A | 2/2 (1) | 0.76 / 1.87 |
| Liu | 9 | 100 | 0 | 2 | 2 | 47 | 26.1 | 2 | A | 2/2 (0) | 0.95 / 1.02 |
| Lucarelli | 84 | 1755 | 0 | 10 | 2 | 60 | 43.585$^*$ | 1 | C | 10/10 (4) | 1.66 / 3.34 |
| Okuonghae | 16 | 92 | 0 | 2 | 2 | 75 | 4.1 | 2 | A | 1/2 (0) | 4.7e4 / 1.2e5 |
| Oliveira | 12 | 120 | 0 | 2 | 2 | 134 | 12.9 | 2 | A | 1/2 (0) | 912 / 8.1e4 |
| Perelson | 3 | 16 | 0 | 2 | 1 | 0 | 0.0 | 3 | A | 0/2 (0) | <0.01 / <0.01 |
| Rahman | 9 | 23 | 0 | 2 | 2 | 70 | 3.7 | 2 | A | 2/2 (2) | <0.01 / 0.03 |
| Raia | 39 | 205 | 5 | 2 | 2 | 32 | 3.3 | 2 | A | 2/2 (0) | 8.22 / 5.10 |
| Raimundez | 136 | 627 | 79 | 2 | 2 | 64 | 18.6$^*$ | 1 | A | 2/2 (0) | 1.17 / 1.43 |
| SalazarCavazos | 6 | 18 | 0 | 2 | 2 | 41 | 18.7 | 2 | A | 2/2 (0) | 102 / 124 |
| Schwen | 30 | 286 | 15 | 2 | 2 | 42 | 17.1 | 2 | A | 2/2 (0) | 8.20 / 8.40 |
| Smith | 25 | 62 | 9 | 2 | 2 | 32 | 7.8 | 2 | A | 1/2 (0) | 1.5e4 / 1.5e4 |
| Sneyd | 15 | 135 | 0 | 2 | 2 | 56 | 38.1 | 1 | A | 2/2 (2) | 2.22 / 2.56 |
| Weber | 36 | 135 | 5 | 2 | 2 | 93 | 4.1 | 1 | A | 2/2 (0) | 1.00 / 15.95 |
| Zhao | 28 | 82 | 0 | 3 | 2 | 70 | 29.6 | 2 | A | 3/3 (1) | 0.35 / 0.35 |
| Zheng | 46 | 60 | 0 | 2 | 2 | 95 | 30.0 | 1 | A | 2/2 (2) | 1.58 / 20.82 |

they determine the few selected parameters. Identifiability agreed with the profiled information $I_j = 1/[(A^\top A)^{-1}]_{jj}$ in every evaluable problem (Methods, *Identifiability of the selected parameters, Formal condition*).

## Prediction of held-out experiments

We fitted the reduced model ($\theta_S$, $\theta_C$) and the full model to training data from $\theta_0$, splitting by experimental condition (21 problems) or by time (9; Blasi could not be split), and evaluated both on the held-out data (Methods, *External validation*). With the non-selected parameters at $\theta_0$, the reduced model predicted at least as accurately (test $\chi^2/n$ within 5% or lower) in 26 of 29 comparable problems (23 of the 26 in Fig. 3a, which omits three problems fitted far below the noise level); the exceptions were Raia, Fiedler and Chen, and where the refitted full model had a much larger test error the difference was large (Brannmark, 6.0 against 6,277).

This comparison favours the reduced model, because $\theta_0$ was estimated from all data, test experiments included (Fig. 3c; Methods, *External validation, Nested validation without the $\theta_0$ anchor*): $\theta_0$ alone predicted better than the reduced model in 20 of 29 problems (median ratio 0.90), and a nested reduced model with the non-selected parameters fixed at a training fit instead of $\theta_0$ predicted worse in 21 of 29 (median ratio 1.35). The advantage on real data reflects anchoring at published estimates, not the selected subset alone.

Without this confound, on synthetic data from a known truth within $\pm 20\%$ of $\theta_0$, with noise of the recorded size and 20 seeds (Methods, *External validation, Synthetic experiment*), which $\theta_0$ has not seen but still locates, so that the test is local, the reduced model reached $e_{\mathrm{pred}} \leq \tau_e$ in 56% of the cases with signal (178 of 320) and predicted at least as well as the full model in 57% (160 of 282). The outcome was stable within problems (Fig. 3b): better in almost every seed where the refitted full model had a much larger test error (Weber 20 of 20), worse in others (Oliveira 0 of 20). We found no overall predictive advantage of the reduced models, and equivalence was not tested. Success decreased with the distance of the truth from $\theta_0$ (64%, 54% and 51% at $\pm 5\%$, $\pm 20\%$ and $\pm 50\%$) and grew with the true change in noise units (Methods, *External validation, True signal*).

## Comparison with existing selection criteria

We compared our subsets with those of pivoted QR in the standard form of Golub, Klema and Stewart[27, 28], Brun's collinearity index[21], Morris screening[15] and profile likelihood[7, 16], each choosing $k = |S|$ biological parameters from the training outputs, and with two null baselines: the calibration parameters alone ($k = 0$) and 20 random $k$-subsets. Every subset went through the same tests: admissibility at $\pm 5\%$ and $\pm 20\%$, prediction of held-out experiments and prediction of synthetic data (Methods, *Comparison with other selection criteria*; Table 2, Fig. 4). Because our main subsets were selected on all outputs, we also repeated our selection on the training outputs alone (the scenarios remain those of our selection; Methods, *Comparison with other selection criteria, Selection on the training outputs and other families of truths*).

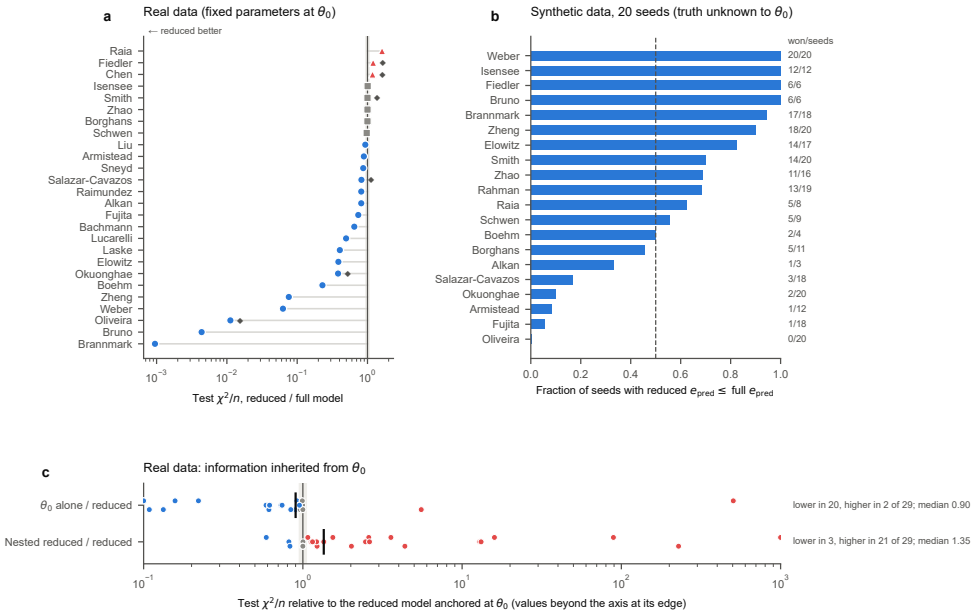



**Fig. 3 Prediction of held-out experiments by the reduced and the full model. a**, Real data: ratio of the test $\chi^2/n$ (mean squared noise-scaled residual on the held-out experiments) of the reduced model to that of the full model re-estimated on the same training data, on a logarithmic scale ($n = 26$ problems), with the non-selected parameters kept at $\theta_0$. Left of the vertical line (circles) the reduced model predicts better; the grey band marks equivalence within 5% (squares); triangles mark worse predictions; ♦ flags a reduced test $\chi^2/n$ above 10, where a favourable ratio does not imply a good absolute fit. **b**, Synthetic data from a known truth, which $\theta_0$ has not seen: for each problem, fraction of the 20 seeds in which the reduced model predicted at least as well as the full model (smaller $e_{\mathrm{pred}}$), among the seeds in which the parameter change exceeded the noise (seeds won / seeds with signal on the right; problems with fewer than three such seeds are not shown). Problems are ordered by this fraction. **c**, Real data, information inherited from $\theta_0$: test $\chi^2/n$ of the published estimate $\theta_0$ without re-estimation, and of the nested reduced model (non-selected parameters fixed at a full-model fit to the training data instead of $\theta_0$), each relative to the reduced model anchored at $\theta_0$; one point per problem ($n = 29$), vertical bar at the median; blue, lower by more than 5%; red, higher by more than 5%; grey band, within 5%. Problems without a training/test split (Bertozzi, Blasi), without a full re-estimation (in **b**, Chen and Raimundez) or fitted far below the noise level (**a**: Giordano, Perelson, Rahman) are not shown. Per-problem values and scatter plots are given in Supplementary Notes *Prediction of held-out experiments* and *Nested validation without the $\theta_0$ anchor*.

On real data no criterion differed significantly from ours (Wilcoxon $P = 0.16$–$0.46$; Fig. 4c; Methods, *Statistics*) except calibration-only re-estimation, which was slightly better (median ratio 0.99, $P = 0.003$): in sloppy models[11] many small subsets predict similarly when the others stay at $\theta_0$.

The criteria differed in reproducing the full model (Fig. 4a). At ±5%, our subsets were admissible in all 28 comparable problems, Brun's in 26, Morris's in 23, profile likelihood's in 18 of 27 and pivoted QR's in 17 (Holm-adjusted McNemar $P = 0.5$, 0.13, 0.012 and 0.004). This favours our method, which selects by this very test; selected on the training outputs, our subsets passed in 24 of 27. The null baselines show that the test is often not selective: the calibration parameters alone passed in 9 of 28 problems and random subsets in a median of 16 (14–19 over the 20 draws), so the test separated subsets in about 12 problems. Part of this reflects the scenarios, which also perturb the calibration parameters that are then re-estimated (a median of 30% of the output change in the 14 problems that have them, up to 93% in Laske). With these parameters kept at $\theta_0$ and the same biological perturbations, our subsets remained admissible in 13 of the 14 (Bachmann rose to 44.1%), calibration alone fell from 9 to 6 and the admissible random draws from 220 to 192 of 280; in Laske only 2 of 20 random pairs still passed, whereas ours did (Methods, *Comparison with other selection criteria*, *Calibration parameters fixed*; Supplementary Note *Admissibility with the calibration parameters fixed*). At ±20%, with the same directions rescaled, the counts were 27, 25, 23, 18 of 27 and 16 ($P = 0.5$, 0.25, 0.023 and 0.004).

On synthetic data (Fig. 4b; 281 cases with signal in 24 problems), our subsets reached $e_{\text{pred}} \leq \tau_e$ in 56% of the cases, Brun's in 55%, Morris's in 48%, profile likelihood's in 39% and pivoted QR's in 30%; resampling problems, the advantage was 8 percentage points over Morris (95% confidence interval 2–16; Holm-adjusted $P = 0.010$), 17 over profile likelihood (7–28; $P = 0.002$), 26 over pivoted QR (12–40; $P < 0.001$) and 1 over Brun ($-1$ to $+3$; $P = 0.43$). As these truths belong to the family of perturbations of the admissibility test, we repeated the comparison with log-uniform and sparse truths (Methods, *Comparison with other selection criteria*, *Selection on the training outputs and other families of truths*): the advantage over profile likelihood and pivoted QR persisted (log-uniform 50% against 37% and 24%; sparse 54% against 40% and 39%; $P \leq 0.019$), that over Morris only for sparse truths, and no family separated our method from Brun's. Selected on the training outputs, our subsets reached 51%, 48% and 55% in the three families, more often than profile likelihood and pivoted QR ($P \leq 0.046$) and without a significant difference from Brun's (equivalence was not tested). The median prediction error differed significantly only from that of pivoted QR (Supplementary Note *Comparison of selection criteria on synthetic data*). What distinguishes our method is thus not its ranking, which it shares with Brun's index and orthogonalization methods[23, 25], but the verification that follows it.

## Identifiable and sufficient

Profiling the parameters selected by every criterion with the same procedure (Methods, *Identifiability of the selected parameters*, *Subsets of other criteria*), 86% of ours were practically identifiable over the 27 common problems, against 85% for

**Table 2 Comparison of selection criteria and null baselines.** All criteria choose the same number $k = |S|$ of biological parameters (except our subsets selected on the training outputs, whose size the method chooses), from the training outputs except for our main subsets (selected on all outputs), and are evaluated with the same protocol and the fixed scenario set. Admissible: problems in which re-estimating the subset reproduces the full model (median $e_{\mathrm{fit}} \leq \tau_e$) for parameter changes of ±5% and ±20%; for random subsets, median over the 20 draws; for calibration only, without the rule $|S| \geq 2$. Identifiable: selected parameters whose profile likelihood in the reduced problem is bounded on both sides, and in parentheses the problems in which every selected parameter is identifiable and the subset is admissible at ±5%, over the 27 problems common to all criteria (26 for our subsets selected on the training outputs). Synthetic: fraction of the cases with signal in which $e_{\mathrm{pred}} \leq \tau_e$, for truths uniform within ±20% of $\theta_0$ (281 cases common to all criteria, 24 problems, 20 seeds; difference from our method with a 95% confidence interval from resampling problems), log-uniform within a factor of two (217 cases, 25 problems) and sparse (147 cases, 24 problems), ten seeds each. Real data: median ratio of the test $\chi^2/n$ to that of our method (below 1, better than ours) and Wilcoxon $P$ over problems (unadjusted; exploratory). Ours (all), (train.): our subsets selected on all outputs and on the training outputs; Profile lik., profile likelihood; Calib., calibration only. Pivoted QR is the standard Golub–Klema–Stewart subset selection; the earlier variant is reported in the Supplementary Information. Per-problem values are given in the Supplementary Notes on the comparison of selection criteria.

| | Admissible | | Identifiable | Synthetic: $e_{\mathrm{pred}} \leq \tau_e$ | | | | Real data |
|---|---|---|---|---|---|---|---|---|
| | | | | Ours − other | | | | |
| Criterion | ±5% | ±20% | Param. (both) | Uniform | (95% CI) | Log-unif. | Sparse | $\chi^2$ ratio ($P$) |
| Ours (all) | **28/28** | **27/28** | 57/66 (**21**) | 56% | – | 50% | 54% | – |
| Ours (train.) | 24/27 | 23/27 | 50/61 (16) | 51% | +5 [0, 12] | 48% | 55% | 1.00 (0.17) |
| Brun | 26/28 | 25/28 | 56/66 (18) | 55% | +1 [−1, 3] | 50% | 52% | 1.00 (0.27) |
| Morris | 23/28 | 23/28 | 53/66 (14) | 48% | +8 [2, 16] | 47% | 46% | 1.00 (0.46) |
| Profile lik. | 18/27 | 18/27 | 57/66 (14) | 39% | +17 [7, 28] | 37% | 40% | 1.00 (0.44) |
| Pivoted QR | 17/28 | 16/28 | 50/66 (7) | 30% | +26 [12, 40] | 24% | 39% | 1.00 (0.16) |
| Random (20) | 16/28 | 16/28 | – | – | – | – | – | 1.00 (0.27) |
| Calib. ($k = 0$) | 9/28 | 9/28 | – | – | – | – | – | 0.99 (0.003) |

Brun, 86% for profile likelihood, 80% for Morris and 76% for pivoted QR (Table 2), so identifiability alone separates only the criteria without a redundancy test. Requiring every selected parameter to be identifiable *and* the subset admissible gave 21 of 27 problems for our subsets, 18 for Brun's, 14 for Morris and for profile likelihood, and 7 for pivoted QR (Holm-adjusted McNemar $P = 0.38$, 0.047, 0.078 and 0.002; Supplementary Note *Identifiability of the subsets of other criteria*); profile likelihood selects identifiable parameters that often do not suffice.

## Robustness to the size of parameter changes

Admissibility is decided at ±5%. Re-evaluating each final subset at ±1 to ±50% (15 scenarios per level), we classified the median $e_{\mathrm{fit}}$ into patterns A (admissible at all levels), B (failing first at ±20% or more), C (irregular) and D (failing at ±10% or less; Methods, *Robustness patterns*). With the same directions rescaled, 26 of the 31 problems follow pattern A (Fig. 5), four pattern B (Armistead, Bertozzi, Blasi, Elowitz, already at 33–40% at ±5%), one pattern C (Lucarelli) and none D; with independent directions, 24 follow A, one B and six C, and 22 follow A in both

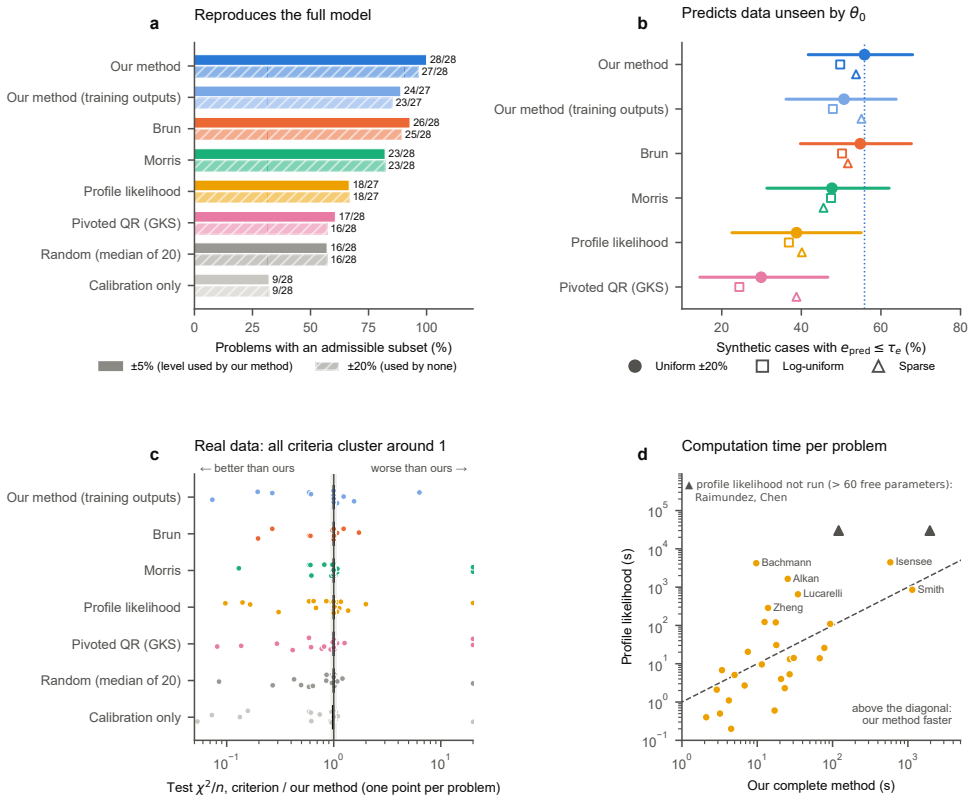



**Fig. 4 Our method compared with four existing selection criteria and null baselines. a**, Number of problems in which the subset of each criterion is admissible, i.e. its re-estimation reproduces the full model, for parameter changes of ±5% (the level used by our method) and ±20% (used by none); 28 problems (27 for profile likelihood and for our subsets selected on the training outputs); random subsets, median over 20 draws. **b**, Fraction of synthetic cases with signal in which the subset predicts the truth within $\tau_e$, for three families of truths (filled circles, uniform within ±20% of $\theta_0$, with 95% confidence intervals from resampling problems; squares, log-uniform; triangles, sparse); dotted line, our method with uniform truths. **c**, Real data: ratio of the test $\chi^2/n$ of each criterion to that of our method, one point per problem, vertical bar at the median (log scale; grey band, ±5%; values outside the axis are drawn at its edge); all criteria cluster around 1. **d**, Computation time of our method (selection, admissibility test and backward elimination; the nonlinear check of the final subset in the six problems selected to first order is not included) and of profile likelihood per problem (log scales); points above the diagonal are faster with our method; problems in which profile likelihood was not run because they have more than 60 free parameters (Chen and Raimundez) are shown at the top.

designs (five evaluated to first order beyond ±5%; Supplementary Note *Robustness to the size of parameter changes*). At ±5% the median pass rate is 93%, and 13 problems reproduce the full model in every valid scenario.

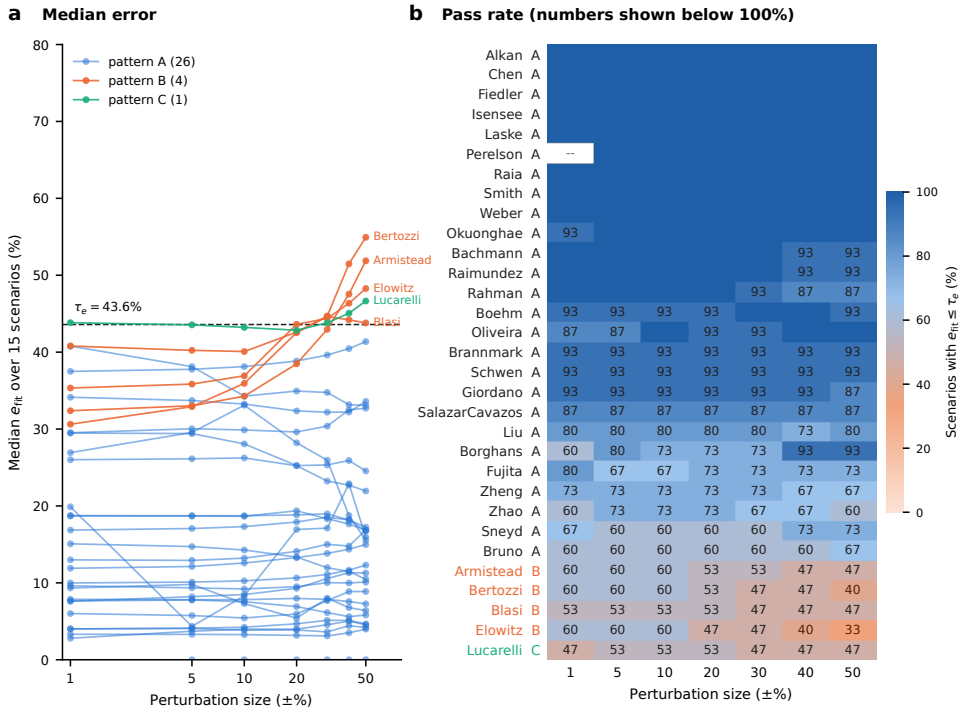



**Fig. 5 Robustness of the selected subsets to the size of parameter changes. a**, Median relative error after re-estimation ($e_{\mathrm{fit}}$) over 15 random scenarios per level as a function of the perturbation size, for the final subset of each of the 31 admissible problems, coloured by robustness pattern (A, admissible at all levels; B, first failing at ±20% or more; C, irregular); dashed line, $\tau_e = 43.6\%$. **b**, Pass rate: percentage of the scenarios with $e_{\mathrm{fit}} \leq \tau_e$, with the value written when it is below 100%. The same fifteen random directions are used at every level. –, no valid scenario (Perelson at ±1%, where the output change is below the numerical noise floor). For the six problems selected to first order (marked in Table 1: Alkan, Bachmann, Chen, Isensee, Lucarelli, Raimundez) the values beyond ±5% are first-order $e_{\mathrm{fit}}$.

## Sensitivity to method settings and computational cost

Re-running the complete method on 26 problems with one setting changed at a time (Methods, *Settings*), the 25 admissible subsets were unchanged by the finite-difference step, the integrator tolerances and $\tau_R$; the perturbation size changed one subset and the conditioning thresholds three or four, without changing the number of admissible problems. With $\tau_e = 0.10$, 0.20, 0.30 and 0.50, 10, 15, 18 and 23 subsets were unchanged, and $\tau_e = 0.10$ needs a mean of 3.7 biological parameters: the threshold is an explicit accuracy–parsimony trade-off, not a tolerance validated for any application. In noise units, the median remaining error per measurement was below $0.1\sigma$ in 17 of 31 problems and below $\sigma$ in 29 (Supplementary Table S34). The random seed and the number of scenarios were not varied.

The sensitivities need $3p$ independent simulations and the admissibility test 15 re-estimations per candidate subset. With the sensitivities in parallel, the selection took a few seconds to 33 min per problem (median 19 s); profile likelihood was faster on small problems (its first round in 16 of 28) but much slower on large ones (Bachmann 70 min on 60 cores against 10 s; Isensee 5.2 h against 10 min, plus 3.9 h for the nonlinear check of our subset; Methods, *Computational cost*). The method itself cost about 23 h of processor time for the 34 problems (1.3 h for the selection and 21.3 h for the nonlinear check of the six first-order problems); the remaining ∼674 h of the 697 h of all analyses were spent on the validations and comparisons of this paper (Supplementary Note *Computing time of the analyses*).

## Formal verification of the mathematical foundations

The statements on which the method relies are verified in Lean 4 with Mathlib (32 modules, 309 theorems and lemmas, and 28 benchmark models), without unproved assumptions (`sorry`) or additional axioms (Methods, *Formal verification*; Supplementary Tables S4 and S7). Table 3 states, for each step, what is proved, when it applies to a numerical result and to which results of this study. Lean does not verify that the optimiser converges, the floating-point values ($\kappa$, VIF, $\cos\Delta$, $e_{\mathrm{fit}}$, $I_j$), the identifiability of the full model, or the translation beyond the nominal parameters and base condition; these are checked numerically. Two statements are central:

**Theorem 1.** *For vectors $d \neq 0$ and $b \neq 0$ in a real inner-product space,*

$$\min_{\lambda \in \mathbb{R}} \frac{\|d - \lambda b\|}{\|d\|} = \sqrt{1 - \cos^2(d, b)}.$$

**Theorem 2.** *Let $\mathcal{A}$ be any admissibility predicate on finite sets and $S$ an admissible set. Backward elimination (removing one element at a time whenever the remaining set is still admissible, until no element can be removed) returns an admissible subset of $S$ from which no element can be removed without losing admissibility. If, in addition, $\mathcal{A}$ is monotone (every superset of an admissible set is admissible), no proper subset of the result is admissible.*

Theorem 1 links $\tau_e$ to the cosine: if $\cos\Delta \geq c$, rescaling leaves a relative error of at most $\sqrt{1 - c^2}$; admissibility does not require $\cos\Delta \geq c$ and is decided by the computed $e_{\mathrm{fit}}$. The first part of Theorem 2 holds for the predicate as implemented, so the subsets are minimal relative to the implemented test, including $|S| \geq 2$, not of minimal cardinality; the second part requires monotonicity, which held in all 660 nested pairs evaluated during the selections (Methods, *Selection*, *Monotonicity log*). Further statements cover conditioning, re-estimation, transformed outputs and the first-order approximation (Methods, *Formal verification*, *Further verified statements*). A sufficient condition for practical identifiability bounds the profile interval through $I_j$, the residual gradient and a bound $\varepsilon$ on the nonlinear remainder (Methods, *Identifiability of the selected parameters*, *Formal condition*); its linear version agreed with the profiles for all 72 evaluable parameters, and with the

remainder estimated at the profile points its hypotheses held for 24 parameters, never for a non-identifiable one (a numerical check, not a certificate). Over the admissible set they held for 9 of the 74 parameters within a factor of 1.5 of the optimum, 7 within a factor of 2 and none over the whole range (Methods, *Identifiability of the selected parameters*, *Hypotheses over the admissible set*), so identifiability over the range of the study rests on the computed profiles. A conclusion of Table 3 applies to a result only when the model hypotheses are verified in Lean and the computed quantities satisfy the stated inequalities (Supplementary Table S5).

# Discussion

In 31 of 34 benchmark problems with real data, re-estimating a subset of usually two biological parameters (the imposed minimum; one sufficed in 15 of 27) with the calibration parameters reproduced, within the chosen tolerance, the local response of the measured outputs, and with the other parameters at their published values the data determined the selected parameters in 25 problems (23 with a misfit-scaled threshold); the number of biological parameters to estimate fell from 750 to 74 (median 15 to 2; Table 1). This is a parameter subset selection[28] checked twice, for sufficiency and for identifiability. With all parameters free most selected parameters are not identifiable, the sloppy-model picture[11, 12]; our statement is more concrete because it refers to the outputs actually measured and is checked by re-estimation and profiles rather than inferred from the sensitivity spectrum. It is not a claim of predictive superiority: on data unseen by the published estimate the reduced models predicted neither consistently better nor worse, and their advantage on real data reflects the published values.

With the other parameters at their published values, all criteria predict held-out data similarly (equivalence was not tested), and those with a redundancy test choose similarly identifiable parameters. They differ in whether their subsets reproduce the full model, but this test is selective only in some problems: part of the success of the null baselines reflects the perturbation of the calibration parameters, and even without it almost any subset passed in six problems (Brannmark, Fujita, Isensee, Raia, Raimundez, Smith), so there the test does not single out the selected parameters. Our subsets were identifiable and sufficient in 21 of 27 problems, against 18 for Brun's index ($P = 0.38$) and 14 for profile likelihood ($P = 0.078$), a count that partly reflects selection by the sufficiency test itself (16 of 26 with selection on the training outputs). The ranking is not new—$E_j$ is the diagonal of the Fisher information matrix of the scaled outputs and, up to a constant, Brun's squared importance measure[21], and orthogonalization methods[23–25] combine it with a greedy redundancy test, much like stage 1—and energy alone sufficed in only seven problems. What the method adds is the verification by re-estimation on the outputs of interest, backward elimination, and failures that become visible (Beer, Lang, Crauste) instead of untested subsets.

Other approaches answer different questions: likelihood-based observability analysis[9] assesses prediction uncertainty after calibration, Bayesian approaches[39] need many more simulations, structural identifiability analysis[2–4] addresses

**Table 3 What the formal verification guarantees.** For each step of the method, the modules of the Lean development (numbers of Supplementary Table S4), the statement proved for any model that satisfies the stated hypotheses, the condition under which it applies to a numerical result, and the results of this study to which it applies. Levels of guarantee and identifiability per problem are given in Table 1; the formal statements are listed in Supplementary Table S7, the scope of each guarantee in Supplementary Table S5, and the correspondence between each claim of this paper and the Lean theorem that supports it in Supplementary Table S43.

| Step of the method (modules) | Statement proved in Lean | Applies to a result when | In this study |
|---|---|---|---|
| Ranking by energy (1, 2, 11, 13) | The first-order output error when parameters are fixed is bounded by the discarded sensitivity energy; the bound transfers to transformed outputs | $R_{\mathrm{var}}(S) \geq \tau_R$ and the model is verified | None: Alkan and Perelson satisfy $R_{\mathrm{var}} \geq \tau_R$, but their models are not verified |
| Selection with $\kappa$ and VIF (7–9) | The greedy subset satisfies $\kappa \leq \tau_\kappa$ and $\mathrm{VIF} \leq \tau_{\mathrm{VIF}}$; discarding a rejected candidate permanently is correct; the linearised least-squares estimates are unique and stable | The computed $\kappa$ and VIF are below the thresholds (unit-length columns for the least-squares statement) | Selection: all 31 subsets (centred columns); least squares: 30 for $\theta_S$ with $\Theta_C$ fixed (not Rahman); with $\Theta_C$ re-estimated, not satisfied in 10 |
| Admissibility threshold (3, 4, 20–22) | Theorem 1: rescaling an output change with cosine $c$ leaves a relative error $\sqrt{1-c^2} = \tau_e$ | Median $\cos\Delta \geq c$ | Level 2 or 3: 17 subsets |
| Re-estimation (5, 12) | More re-estimated parameters never increase the fitting error; $e_{\mathrm{fit}} \leq e_{\mathrm{rel}}$; the first-order $e_{\mathrm{fit}}$ differs from the exact one by a second-order term | Exact minimum; the computed $e_{\mathrm{fit}}$ (local optimum) bounds it from above, so admissible as computed implies admissible exactly, not conversely | 25 by nonlinear re-estimation; 6 first order, then checked nonlinearly at $\pm 5\%$; Beer and Lang first order; Crauste not defined |
| Backward elimination (6) | Theorem 2: the result is admissible and no parameter can be removed | Relative to the implemented test, including $\lvert S\rvert \geq 2$ | All 31 subsets (non-trivial only for Bruno, Lucarelli and Zhao) |
| Model hypotheses (14–18, 23–31) | The solution exists on the horizon, stays non-negative and is differentiable in the parameters | Checked by the Lean kernel at the nominal parameters and base condition (first segment of piecewise models) | 28 models (27 of the 34 problems); Crauste only for $T \leq 1$; the selected initial conditions of Bertozzi and Bruno are not parameters of the Lean files |
| Identifiability (32) | If $t_j^*$ is smaller than the distance to both bounds, $\theta_j$ is practically identifiable | $I_j$, $\gamma$ and $\varepsilon$ computed at the profile points satisfy the inequality | Hypotheses satisfied at the evaluated profile points (not a certificate): 24 parameters; all of $S$ in 7 problems |

noise-free data, manifold-based methods[40] rewrite the model in emergent parameter combinations, balanced truncation[41, 42] reduces state variables, and hierarchical optimisation[43] eliminates calibration parameters without selecting among biological ones.

The study has limitations. (i) The analysis is local: sensitivities are computed at $\theta_0$, assumed to be a good fit, admissibility is tested around it, and the synthetic truths are generated around it. (ii) Because the non-selected parameters stay at $\theta_0$, predictions on real held-out data benefit from information that $\theta_0$ already contains. (iii) The minimum subset size of two is a conservative choice: one parameter already reaches the threshold in about half of the problems and is as identifiable and predictive there, but fails in more individual scenarios, especially for large changes (Supplementary Note *Minimum subset size: one parameter or two*). (iv) The error threshold $c = 0.9$ is a design choice relative to the perturbation, and in 25 of the 31 problems a $\pm 5\%$ perturbation changes the outputs by less than the measurement noise (by less than $10^{-4}$ of it in Giordano and Perelson), so admissibility at this level concerns the response of the model rather than differences detectable in the data. (v) The explicit finite-perturbation certificates of the formal development (Methods) are too conservative to certify any scenario at $\pm 5\%$ in our models (constants $10^2$–$10^{24}$), so admissibility rests on the local theorems and on the numerical test. (vi) The method assumes smooth models, Gaussian noise with the recorded standard deviations and re-estimable calibration parameters; in five problems the residuals do not match the recorded noise levels, one problem with 9,169 conditions was too expensive to analyse on its measured outputs, and three problems had no admissible subset.

For a modeller, the practical message is simple: before calibrating a large model, select a small set of biological parameters, verify that re-estimating them with the calibration parameters reproduces the full model (and whether the calibration parameters alone already do) and that their profiles are bounded, keep the others at their published values, and check the reduced model on held-out experiments. The selection is cheap (median 19 s per problem), and the subset is often one of many admissible ones: selected on the training outputs, 20 of the 29 problems chose a different one, mostly another pair. Natural extensions are forward sensitivity equations for genome-scale models[44], several operating points and experimental design[45]. In most of the benchmark problems, recalibrating a published model with new data could thus start from two checked parameters instead of a median of 15.

# Methods

## Models, data and simulation

We used the problems of the PEtab benchmark collection[10, 36] that provide real measurements, an SBML model and nominal parameter values (34 problems; Table 1). This is a secondary analysis of published, de-identified data, for which ethics approval was not required (Ethics declarations). Each problem specifies the observables $g_i$, their transformation $h$ (linear, log or $\log_{10}$), the noise standard deviation $\sigma_i$ of each measurement, the experimental conditions (parameter or initial-value overrides) and, where applicable, pre-equilibration conditions. Models

were simulated with libRoadRunner[46] (CVODE, relative and absolute tolerances $10^{-7}$ and $10^{-10}$). Initial assignments that depend on parameters were re-evaluated for every parameter vector, condition-specific overrides were applied before simulation, pre-equilibration was run to steady state before the main simulation, and models with discontinuous inputs were integrated piecewise between discontinuities. Simulated observables were compared with the reference simulations distributed with the benchmark wherever available (Supplementary Note *Verification of the simulator*, Supplementary Table S10). *Crauste.* To establish why no local reduction exists for this problem, each of its 12 estimated parameters was changed by ±1% one at a time and the states were simulated over the measured horizon of 28 days (`crauste_blowup.py`); the solution grows without bound before the end of the horizon in 9 of these 24 changes and in 7 of the 15 admissibility scenarios at ±5%, with the same counts for relative integrator tolerances from $10^{-10}$ to $10^{-5}$. One further benchmark problem with real data (Froehlich, 4,231 parameters and 9,169 experimental conditions) was analysed only on its state trajectories, because each evaluation of its measured outputs requires about 9,000 simulations; it is not counted in the 34 problems (see Computational cost).

## Measured outputs and parameter classes

For a parameter vector $\theta \in \mathbb{R}^p$ (in the units of the benchmark, excluding noise parameters), the vector of scaled outputs is $y(\theta) \in \mathbb{R}^m$ with entries

$$y_i(\theta) = \frac{h\big(g_i(x(t_i;\theta,c_i),\theta)\big)}{\sigma_i}, \qquad i = 1,\dots,m, \tag{1}$$

where $x(t;\theta,c)$ is the state trajectory under condition $c$ and $(t_i, c_i)$ are the time and condition of measurement $i$. Parameters that appear in the SBML model are biological ($\Theta_B$); parameters that appear only in the observation functions and not in the model or the condition table (scalings and offsets) are calibration parameters ($\Theta_C$). Only biological parameters are candidates for selection; calibration parameters are always re-estimated.

## Sensitivities, energy and conditioning

The relative sensitivity matrix $J \in \mathbb{R}^{m\times p}$ is computed by central differences,

$$J_{:,j} = \frac{y(\theta_0 + \delta\theta_{0,j}e_j) - y(\theta_0 - \delta\theta_{0,j}e_j)}{2\delta}, \tag{2}$$

which approximates $\partial y/\partial \log\theta_j$. The energy of parameter $j$ is $E_j = \|J_{:,j}\|^2 = (J^\top J)_{jj}$ and the fraction of energy captured by a subset $S \subseteq \Theta_B$ is $R_{\mathrm{var}}(S) = \sum_{j\in S} E_j / \sum_{j\in\Theta_B} E_j$. For a subset $S$, the columns of $J_S$ are centred and scaled to unit variance, giving a matrix $F$ whose Gram matrix is proportional to the correlation matrix $R_S$ of the sensitivity columns; the condition number is $\kappa(S) = \sigma_{\max}(F)/\sigma_{\min}(F)$ and the largest variance inflation factor is $\mathrm{VIF}(S) = \max_j [R_S^{-1}]_{jj}$[18, 21]. Scaling removes differences of scale between

parameters, so that $\kappa$ and VIF measure only the angles between sensitivity directions. Because $R_S$ is a principal submatrix of the correlation matrix of all columns, both quantities can only increase when parameters are added. The statements of the formal development on the linearised least-squares problem are formulated for columns scaled to unit length without centring, $Z = J_S \operatorname{diag}(\|J_{:,j}\|)^{-1}$; computed in that form for the final subsets (`conditioning_check.py`), $\kappa(Z) \leq 6.6$ and $\mathrm{VIF}(Z) \leq 10$ in 30 of the 31 subsets, whereas in Rahman $\kappa(Z) = 12.4$ and $\mathrm{VIF}(Z) = 39$, so the formal statement on the linearised problem does not apply to that subset.

## Selection

*Greedy selection.* Candidates are visited in a given order. A candidate $j$ is accepted if $\kappa(S \cup \{j\}) \leq \tau_\kappa$ and $\mathrm{VIF}(S \cup \{j\}) \leq \tau_{\mathrm{VIF}}$; otherwise it is discarded permanently. Because $\kappa$ and VIF can only increase when columns are added, a discarded candidate would also be rejected later. *Stage 1* visits biological parameters in decreasing order of $E_j$ and stops when $R_{\mathrm{var}}(S) \geq \tau_R$. *Stage 2*, used if the stage-1 subset is not admissible, visits biological parameters in decreasing order of the finite-perturbation effect $s_j = \|y(\theta_0 + \eta\theta_{0,j}e_j) - y(\theta_0)\|$ and stops at the first admissible subset. *Backward elimination* visits the parameters of the selected subset and removes each one whose removal leaves the subset admissible (between one and five parameters were removed in five problems; Supplementary Table S1).
*Monotonicity log.* All subsets evaluated during a selection and their admissibility are stored; for every nested pair $S \subset T$ of evaluated subsets we count the violations of monotonicity ($S$ admissible, $T$ not), which bear on the second part of Theorem 2.

## Admissibility and re-estimation

Admissibility is evaluated on $N$ random scenarios $\theta^{(k)} = \theta_0 \odot (1 + \varepsilon u^{(k)})$, with $u^{(k)}$ uniform in $[-1, 1]^p$. For each scenario let $\Delta y_{\mathrm{full}} = y(\theta^{(k)}) - y(\theta_0)$. The relative error without re-estimation is $e_{\mathrm{rel}} = \|\Delta y_{\mathrm{full}} - \Delta y_S\| / \|\Delta y_{\mathrm{full}}\|$, where $\Delta y_S = y(\theta^{(k)}_{S \cup \Theta_C}, \theta_{0,\overline{S \cup \Theta_C}}) - y(\theta_0)$ is the output change when only the parameters in $S \cup \Theta_C$ take their scenario values and all others stay at $\theta_0$, and the cosine similarity is $\cos\Delta = \langle \Delta y_{\mathrm{full}}, \Delta y_S \rangle / (\|\Delta y_{\mathrm{full}}\| \|\Delta y_S\|)$. The relative error after re-estimation is

$$e_{\mathrm{fit}} = \min_{\theta_S, \theta_C} \frac{\|y(\theta_S, \theta_C, \theta_{0,\overline{S}}) - y(\theta^{(k)})\|}{\|\Delta y_{\mathrm{full}}\|}, \tag{3}$$

where $\theta_{0,\overline{S}}$ denotes the nominal values of all other biological parameters. The minimisation is carried out in $\log|\theta|$ within $[|\theta_0|/10,\ 10|\theta_0|]$ by a trust-region reflective least-squares method (SciPy[47], at most $100(|S| + |\Theta_C| + 1)$ function evaluations), from two starting points (the scenario value and $\theta_0$); the error without re-estimation is also included, so that $e_{\mathrm{fit}} \leq e_{\mathrm{rel}}$. Scenarios whose full-model output change is below $\max(10^{-10}\|y(\theta_0)\|, 10^{-6})$, or in which the simulation of the full model fails, are discarded, and the number of valid scenarios is reported; this set does not depend on $S$. A scenario in which the simulation with only $S \cup \Theta_C$ at their scenario values fails is kept and counts as not reproduced ($e_{\mathrm{rel}} = \infty$). Parameters

with a negative nominal value keep their sign in the scenarios and in the re-estimation, which acts on $\log|\theta|$; a parameter with nominal value zero (four in Lang and two in Raimundez) is represented by $10^{-12}$, so its relative perturbations and its column of $J$ are negligible and, like any parameter with a zero or negligible column, it is in practice never selected (Supplementary Note *Random scenarios and large problems*). Supplementary Table S3 reconciles the parameter counts used throughout (estimated, biological, calibration and noise parameters, and the free parameters of each analysis). $S$ is admissible if

$$|S| \geq k_{\min} \quad \text{and} \quad \bar{e}_{\text{fit}}(S) = \underset{k}{\text{median}}\, e_{\text{fit}}^{(k)}(S) \leq \tau_e = \sqrt{1-c^2}, \tag{4}$$

the median being taken over the valid scenarios. The reduction in the number of free parameters reported in Results is $1-(|S|+|\Theta_C|)/p$, and $1-|S|/|\Theta_B|$ among biological parameters. Because $e_{\text{rel}}$ bounds the exact $e_{\text{fit}}$ from above, the implementation also accepts a subset whose median $e_{\text{rel}}$ is at most $\tau_e$, which makes re-estimation unnecessary; with nonlinear re-estimation $e_{\text{fit}} \leq e_{\text{rel}}$ in every scenario, so the rule $\min(\text{median}\, e_{\text{fit}}, \text{median}\, e_{\text{rel}}) \leq \tau_e$ reduces to the median of $e_{\text{fit}}$, and it differs only for the first-order $e_{\text{fit}}$ described below. For each admissible subset we also report a level of guarantee: level 1 when only this condition holds, level 2 when in addition the median $\cos\Delta$ is at least $c$, and level 3 when in addition $R_{\text{var}}(S) \geq \tau_R$; these levels are reported, not required. *First-order $e_{\text{fit}}$ and nonlinear check.* For the eight problems in which the cost of simulation made repeated re-estimation impractical (marked in Table 1), the first-order version $e_{\text{fit,lin}} = \min\{e_{\text{rel}}, \min_c \|\Delta y_{\text{full}} - J_{S\cup\Theta_C} c\|/\|\Delta y_{\text{full}}\|\}$, evaluated scenario by scenario, was used during the selection. This equals the minimum relative error attainable if the model response were exactly linear in parameter changes; for a truly linear model it equals the exact $e_{\text{fit}}$, while for nonlinear models it is an approximation that replaces per-scenario optimisation with a single projection; its absolute error is bounded by a second-order term (Supplementary Table S7), but no bound on the relative error was evaluated for the individual scenarios. The admissibility decision during the selection is therefore made by linear algebra alone. For the six of these problems with an admissible subset, the final subset was afterwards re-estimated nonlinearly in the same 15 scenarios, with the same starting points, bounds and evaluation budget as for the other problems, the scenarios being distributed over processes (`efit_nonlinear.py`); it confirmed admissibility in all six (median $e_{\text{fit}}$, the values reported in Table 1 and Fig. 2: Alkan 0.120, Bachmann 0.338, Chen 0.081, Isensee 0.113, Lucarelli 0.43585 against $\tau_e = 0.43589$, Raimundez 0.186; Supplementary Note *Nonlinear check of the first-order problems*). The columns of $J$ for the calibration parameters are computed together with those of the biological parameters, because the first-order re-estimation uses $J_{S\cup\Theta_C}$.

## Settings

Unless stated otherwise: $\delta = 0.01$; $\eta = 0.1$; $\tau_R = 0.89$; $\tau_\kappa = \tau_{\text{VIF}} = 10$; $N = 15$ scenarios; $\varepsilon = 0.05$; random seed 42; $c = 0.9$, so $\tau_e = 0.436$; $k_{\min} = 2$. The sensitivity analysis changed one setting at a time: $\delta \in \{0.1\%, 5\%\}$; relative/absolute integrator

tolerances $10^{-9}/10^{-12}$ and $10^{-6}/10^{-8}$; $\varepsilon \in \{0.01, 0.10\}$; $\tau_R \in \{0.80, 0.85, 0.95\}$; $\tau_\kappa = \tau_{\mathrm{VIF}} \in \{5, 30\}$; and $\tau_e \in \{0.10, 0.20, 0.30, 0.50\}$ (Supplementary Note *Sensitivity to the method settings*). It was run on the 26 problems that could be re-run completely fifteen times in reasonable time; the other eight (Alkan, Bachmann, Beer, Chen, Isensee, Lang, Lucarelli and Raimundez) are those with the most expensive simulations. To relate $e_{\mathrm{fit}}$ to the precision of the data, the remaining error of each scenario was also expressed per measurement in units of the recorded noise, $e_{\mathrm{fit}}\|\Delta y_{\mathrm{full}}\|/\sqrt{m}$ (the outputs are already divided by $\sigma_i$), and its median over the scenarios was reported. *Single-parameter subsets.* Single-parameter subsets were tested on 27 admissible problems (all except the four with the most expensive simulations, Bachmann, Chen, Isensee and Raimundez): the ten biological parameters with the largest $E_j$ and the ten with the largest $s_j$ were each tested alone, with $\Theta_C$ re-estimated, on the same 15 scenarios at $\pm 5\%$ and without the requirement $|S| \geq 2$; the admissible candidate with the smallest median $e_{\mathrm{fit}}$ was then validated like the selected subset, and, to test the rule $k_{\min} = 2$, put through the robustness, $\pm 20\%$ admissibility, synthetic and identifiability tests of the selected pair (5.7 h of processor time, 27 min of elapsed time on an 88-core server); in the synthetic experiment the selected subset predicted better than the best single parameter in 9 of these 15 problems, whereas on real data the single parameter had the lower test $\chi^2/n$ in 10 of 15 (Supplementary Note *Single-parameter subsets*). In the scan of settings, the perturbation size changed the subset only in Zhao, and admissibility only for Perelson at $\pm 1\%$, where the output change falls below the numerical noise floor. The admissibility test uses the median over $N = 15$ scenarios rather than requiring all of them to pass, because individual scenarios near the boundary of the parameter domain can fail through ODE nonlinearity; seven problems pass with 8–9 of the 15 scenarios. For Lucarelli the median $e_{\mathrm{fit}}$ lies within $10^{-3}$ of $\tau_e$ to first order and within $4 \times 10^{-5}$ after nonlinear re-estimation, so its admissibility is borderline, and its median crosses $\tau_e$ at several perturbation sizes (43–47% from $\pm 1\%$ to $\pm 50\%$; pattern C). To relate the size of a perturbation to the noise, we computed, for each admissible problem, the root mean square per measurement of $\Delta y_{\mathrm{full}}$ at $\pm 5\%$ (in units of $\sigma$); it is below 1 in 25 of the 31 problems and below $10^{-4}$ in Giordano and Perelson (Supplementary Table S34).

## External validation

Measurements were split into training and test sets by experimental condition (every third condition after sorting, or the second of two conditions, used for testing) or, for problems with a single condition, by time (the last third of the distinct measurement times). Parameters that only affect test conditions were excluded from re-estimation. In two problems a split was not possible: in Bertozzi the selected parameters act only in the test conditions, and Blasi has a single condition and a single time point. All real-data numbers come from a single run of each fit; the re-estimation of over-parameterised full models can end in different local optima[48]. In five problems even $\theta_0$ has a test $\chi^2/n > 10$ (Chen, Okuonghae, Oliveira, Salazar-Cavazos, Smith); our simulation agrees with the reference simulation of the benchmark for Chen, Salazar-Cavazos and Smith, whereas no reference is available for Okuonghae and

Oliveira. A large $\chi^2/n$ can reflect recorded noise levels that understate the scatter of the data, a misspecified model or inadequate parameter values, and our analysis does not distinguish between these causes. The reduced model (re-estimating $S \cup \Theta_C$) and the full model (re-estimating all parameters) were fitted to the training data from $\theta_0$ with the same least-squares method and bounds, $\hat{\theta} = \arg\min \sum_{i \in \text{train}} (y_i(\theta) - \tilde{y}_i)^2$ over the free parameters of each model within $[|\theta_0|/10, 10|\theta_0|]$, and evaluated by the mean squared noise-scaled residual $\chi^2/n = \frac{1}{n} \sum_{i \in \text{test}} (y_i(\hat{\theta}) - \tilde{y}_i)^2$, where $\tilde{y}_i$ are the scaled measurements. *Synthetic experiment.* In the synthetic experiment, the truth $\theta^*$ had each estimable parameter drawn uniformly within ±20% of $\theta_0$ (parameters acting only in test conditions stayed at $\theta_0$), data were simulated with Gaussian noise of standard deviation $\sigma_i$, and $e_{\text{pred}} = \|y_{\text{test}}(\hat{\theta}) - y_{\text{test}}(\theta^*)\| / \|y_{\text{test}}(\theta^*) - y_{\text{test}}(\theta_0)\|$; the effect of the parameter change was considered larger than the noise when $\chi^2/n$ of $\theta_0$ on the synthetic test data was at least 2. The experiment was repeated with 20 seeds (42–61), each defining a different truth and noise realisation; the full model was not re-estimated for the two problems with more than 100 free parameters after removing those that act only in the test conditions (Chen and Raimundez; Bachmann has 101 parameters that affect the outputs but 84 free ones, and its full model was re-estimated). The experiment was repeated with truths within ±5% and ±50% of $\theta_0$ (seeds 42–51, ten seeds), and compared with the same ten seeds at ±20%. With these amplitudes, the reduced model reached $e_{\text{pred}} \leq \tau_e$ in 64% and 51% of the cases with signal at ±5% and ±50% (54% at ±20% on the same seeds) and predicted at least as well as the full model in 67% and 47%. *True signal.* The true signal of a case is the root mean square, over the test measurements, of $y_{\text{test}}(\theta^*) - y_{\text{test}}(\theta_0)$, which is in units of $\sigma$ because the outputs are scaled by the noise; it is computed from the truth, without noise, whereas the criterion "larger than the noise" above uses the noisy synthetic data. At ±20%, the reduced model reached $e_{\text{pred}} \leq \tau_e$ in 1%, 24%, 44% and 75% of the cases whose true signal was below $0.3\sigma$, between $0.3\sigma$ and $\sigma$, between $\sigma$ and $3\sigma$, and above $3\sigma$; across seeds, the success rate ranged from 31% to 76% and the fraction of cases predicted at least as well as by the full model from 40% to 73%. Problems whose test $\chi^2/n$ at $\theta_0$ was below 0.01 (Giordano, Perelson, Rahman) reproduce the data to far below the recorded noise level and are not shown in Fig. 3a. Test errors were compared through the ratio of the reduced to the reference $\chi^2/n$; ratios between 0.95 and 1.05 were described as similar (a descriptive band, not a test). In the synthetic experiment, a case is a pair (problem, seed); the reduced model predicted at least as well as the full model when $e^{\text{red}}_{\text{pred}} \leq e^{\text{full}}_{\text{pred}}$, and the success rate is the number of cases with signal and $e_{\text{pred}} \leq \tau_e$ divided by the number of cases with signal. *Nested validation without the $\theta_0$ anchor.* For the nested validation without the $\theta_0$ anchor (a partial control of leakage) (Supplementary Note *Nested validation without the $\theta_0$ anchor*), the full model was first re-estimated on the training data (yielding $\hat{\theta}_{\text{train}}$), and the reduced model was then re-estimated with the non-selected parameters fixed at $\hat{\theta}_{\text{train}}$ and search bounds derived from $\hat{\theta}_{\text{train}}$, for 29 problems (Bertozzi and Blasi excluded; for Chen the columns of the finite-difference Jacobian were evaluated in parallel and the integrator tolerances were $10^{-6}/10^{-8}$).

## Comparison with other selection criteria

Each criterion selected the same number $k = |S|$ of biological parameters as our method, from the training data, and the calibration parameters of the problem were always added. *Pivoted QR*: the standard subset selection of Golub, Klema and Stewart[27, 29], as used for parameter selection by Vélez-Reyes and Verghese[28]: the sensitivity matrix of the training outputs at $\theta_0$, with columns normalised to unit length, is decomposed as $U\Sigma V^\top$, and the first $k$ pivots of the column-pivoted QR decomposition of $V_k^\top$ (the $k$ dominant right singular vectors) are selected. A simpler variant, the first $k$ pivots of the column-pivoted QR decomposition of the normalised matrix itself, in which the first pivot is decided by rounding and column order when several columns have nearly equal norm, is reported in Supplementary Note *Pivoted QR selection*. *Brun*[21]: dimensionless sensitivities $s_{ij} = \theta_{0,j}\,\partial y_i/\partial\theta_j$ on the $n$ training outputs (forward differences, relative step $10^{-4}$), importance $\delta_j^{\mathrm{msqr}} = \|s_{:,j}\|/\sqrt{n}$, collinearity index $\gamma_K = 1/\sqrt{\lambda_{\min}(\tilde{s}_K^\top \tilde{s}_K)}$, where $\tilde{s}_K$ is $s_K$ with its columns scaled to unit length, and determinant measure $\rho_K = \det(s_K^\top s_K)^{1/(2k)}$; among all $k$-subsets $K$ of the 15 parameters with the largest $\delta_j^{\mathrm{msqr}}$, the one with $\gamma_K \le 10$ and the largest $\rho_K$ is selected (the smallest $\gamma_K$ if none satisfies the bound). *Morris*[15]: each biological parameter is represented by $x_j \in [0,1]$, mapped linearly onto $\log|\theta_j| \in [\log|\theta_{0,j}| - \log 10,\ \log|\theta_{0,j}| + \log 10]$; ten trajectories (seed 42) on the grid $\{0, \frac{1}{3}, \frac{2}{3}, 1\}$ change one parameter at a time, in random order, by $\Delta = \frac{2}{3}$, giving elementary effects $\mathrm{EE}_j = \|y(x + \Delta e_j) - y(x)\|/\Delta$ on the training outputs, and the $k$ parameters with the largest mean effect $\mu_j^*$ over the trajectories are selected. *Profile likelihood*[7, 16]: the full model was first re-estimated on the training data; each biological parameter was then fixed at five values on each side of the optimum (up to a factor of 10, within the same bounds as all other fits) and all other parameters were re-estimated; whenever a point of a profile improved the optimum by more than one unit of $\chi^2$, the full model was re-estimated from that point and the profiles were repeated (at most three times; ten for Raia). A parameter was classified as identifiable when its profile exceeded $\Delta\chi^2 = 3.84$ (95%, one degree of freedom) on both sides, and the $k$ parameters with the narrowest confidence intervals were selected, the width being $w_j = t_j^- + t_j^+$, where $t_j^\pm$ is the displacement in $\log|\theta_j|$ at which $\Delta\chi^2$ first exceeds 3.84 on each side, interpolated linearly between profile points ($\infty$ if it is not exceeded); ties are broken in favour of the larger of the smaller maxima of $\Delta\chi^2$ on the two sides. It was run for the 26 admissible problems with at most 60 free parameters and, in a separate run with the profile points distributed over 60 cores, for Bachmann and Lucarelli (90 and 72 free parameters); it was not run for Chen and Raimundez because of its cost. Extrapolating from Bachmann with the simulation time and the numbers of free parameters and profile points, Raimundez (126 free parameters, about 570 profile points, 0.3 s per simulation) would need of the order of $10^3$ h of computing time and Chen (155 free parameters, about 1,520 points, 35 s per simulation) of the order of $10^5$ h (Supplementary Note *Profile likelihood*). The criteria were implemented as they are usually applied and therefore differ in some numerical settings (data used to select, derivative scheme and step, candidate set, normalisation and computational budget); these differences are listed in Supplementary Note *Brun*

*and Morris selections.* The *calibration-only* baseline re-estimated $\Theta_C$ alone ($k = 0$), and the *random* baseline used 20 random $k$-subsets of the biological parameters (seeds 0–19); for the latter, the number of admissible problems was counted for each draw and its median over the draws is reported. Every subset was then evaluated with the external validation described above (for random subsets, the median test $\chi^2/n$ over the draws), with the admissibility test of our method (the same 15 scenarios and the same rules, with the fixed scenario set described in Admissibility and re-estimation, at $\varepsilon = 0.05$ and, with the same directions $u^{(k)}$, at $\varepsilon = 0.20$; for $k = 0$ the rule $|S| \geq 2$ was not applied) and with the synthetic experiment (20 seeds). *Selection on the training outputs and other families of truths.* Our main subsets were selected on all measured outputs at $\theta_0$, whereas the other criteria selected on the training outputs; to remove this difference, our selection was repeated on the rows of the training set only (`REPRO_OUTPUTS=train`), and the resulting subsets were evaluated with the same three tests. The synthetic truths of the main experiment, $\theta^* = \theta_0 \odot (1 + 0.2u)$, belong to the same family of perturbations that our method uses to decide admissibility; the comparison of criteria was therefore repeated, with ten seeds (42–51), for two other families: log-uniform truths, $\theta_j^* = \theta_{0,j}\, 10^{v_j}$ with $v_j$ uniform in $[-0.3, 0.3]$, and sparse truths, in which a random 15% of the estimable parameters change by $\pm 50\%$ and the others stay at $\theta_0$. *Calibration parameters fixed.* Because the scenarios perturb every parameter and $\Theta_C$ is always re-estimated, part of the change to be reproduced is produced by $\Theta_C$ itself. For the 14 problems of the comparison that have calibration parameters, the admissibility test at $\pm 5\%$ was therefore repeated with $\Theta_C$ kept at $\theta_0$ (the same 15 scenarios with the components of $\Theta_C$ reset, so the biological perturbations are identical and the comparison is paired) for all subsets: ours, the training-output ones, the other criteria, calibration only and the 20 random draws (378 tasks, 16.7 h of processor time, 34 min on 90 cores). For our subsets the median fraction $\|\Delta y_C\|^2/\|\Delta y_{\text{full}}\|^2$ of the standard scenarios, where $\Delta y_C$ is the output change produced by the calibration parameters alone, was also computed. In the other 14 problems the test is unchanged. The problems entering each comparison follow from these requirements. Admissibility was compared in 28 of the 31 admissible problems: Bertozzi was excluded (see Statistics), Blasi has no training/test split on which the other criteria could select, and Chen was not included because the selections of pivoted QR, Brun and Morris on its training outputs require 1,734 simulations of 35 s each (about 17 h of computing time) before any re-estimation. Profile likelihood was compared in 27 of them (all except Raimundez). The synthetic comparison used the 24 problems with at least one seed in which all five criteria had a result and the parameter change exceeded the noise: of the 28, Raimundez lacks the profile-likelihood subset, and Giordano, Laske and Lucarelli have no seed with signal. The computation time of the complete method was measured by running it from scratch, without cached results, with the sensitivity computations distributed over 16–64 cores; all timings were obtained on a 128-core server shared with other computations and may therefore be inflated by resource contention.

## Statistics

Comparisons between criteria were made over problems. For real data and for the synthetic prediction error, the statistic of each problem was the logarithm of the ratio of the errors of two criteria (for synthetic data, its median over seeds) and the two-sided Wilcoxon signed-rank test was applied. For the fraction of synthetic cases reaching $e_{\mathrm{pred}} \leq \tau_e$, the success rate of a criterion is $\hat{\pi} = \sum_s a_s / \sum_s n_s$, where $n_s$ is the number of cases with signal of problem $s$ that are common to all criteria and $a_s$ the number of them with $e_{\mathrm{pred}} \leq \tau_e$; 95% confidence intervals of the difference $\hat{\pi}_{\mathrm{ours}} - \hat{\pi}_m$ are the 2.5 and 97.5 percentiles of the difference recomputed in 10,000 resamples of whole problems with replacement (seed 0), so that seeds of the same problem are not treated as independent. The two-sided $p$ value of the bootstrap is twice the smaller of the fractions of resampled differences at or below and at or above zero; in the Wilcoxon test, problems with equal errors were discarded (zero method of Wilcoxon). Admissibility, one outcome per problem, was compared with the exact McNemar test: if $b$ problems are admissible only with our subset and $c$ only with the other, $P = \min\{1,\, 2\sum_{i=0}^{\min(b,c)} \binom{b+c}{i} 2^{-(b+c)}\}$. The comparison of admissibility and of the synthetic success rate with the four existing criteria is the primary analysis; within each family of four comparisons (admissibility at one perturbation level, or synthetic success rate) $p$ values were adjusted with the Holm procedure, $\tilde{p}_{(r)} = \max_{q \leq r} \min\{1, (5-q)\, p_{(q)}\}$ for the ordered values $p_{(1)} \leq \cdots \leq p_{(4)}$, and the comparisons on real data are exploratory. The absence of a significant difference is not evidence of equivalence; no equivalence margin was tested. Bertozzi was excluded from the comparison of admissibility because its selected parameters act only in the test conditions; selected on the training outputs, our method finds no admissible subset of two parameters there (a single parameter, `RO_CA`, has $e_{\mathrm{fit}} = 0.08$). Exact $P$ values are reported; for the bootstrap the resolution is $2 \times 10^{-4}$.

## Identifiability of the selected parameters

For each problem with an admissible subset, the reduced problem ($S \cup \Theta_C$ free, other parameters at $\theta_0$) was fitted to all measurements from $\theta_0$, with the method and bounds of the external validation. For each selected parameter $\theta_j$, the profile

$$\Delta\chi^2_j(t) = \min_{\log|\theta_j| = \log|\hat{\theta}_j| + t} \chi^2(\theta) - \chi^2(\hat{\theta}), \qquad \chi^2(\theta) = \sum_i \big(y_i(\theta) - \tilde{y}_i\big)^2, \tag{5}$$

where $\hat{\theta}$ is the optimum and the minimum is taken over the other free parameters of the reduced problem, was computed at ten displacements $t$ per side, with $|t|$ from 0.05 to 4.6; points outside the bounds $[|\theta_0|/10, 10|\theta_0|]$ of $\theta_j$ were not used, and if a point improved the optimum the reduced problem was re-fitted from it and the profile repeated (at most three rounds). $\theta_j$ was classified as identifiable if $\Delta\chi^2$ exceeded 3.84 on both sides, practically non-identifiable if on one side, and flat otherwise[7]; a robust variant scaled the threshold by $\max(1, \chi^2_{\mathrm{min}}/\mathrm{d.f.})$. *Formal condition.* Let $r(\delta)$ be the vector of noise-scaled residuals of all measurements when the logarithms of the free parameters move by $\delta$ from the optimum, $r_0 = r(0)$, so that

$\chi^2 = \|r\|^2$, and let $A$ be its sensitivity matrix at the optimum (central differences, step 0.01 in $\log|\theta|$). The profiled information $I_j$ was computed as the squared norm of the component of the column of $A$ for $\theta_j$ orthogonal to the other columns, which equals $1/[(A^\top A)^{-1}]_{jj}$ when $A^\top A$ is invertible; it is the output change, in noise units, that a unit change of $\log|\theta_j|$ produces when the other free parameters compensate as well as possible. The residual gradient $\gamma$ is the norm of the projection of $r_0$ onto the range of $A$ (zero at an exact optimum), and $\varepsilon$ is the largest nonlinear remainder $\|r(\delta) - r_0 - A\delta\|$ over the computed profile points. The Lean theorem `practically_identifiable_grad` states that if $|\langle r_0, A\delta\rangle| \le \gamma\|A\delta\|$, the nonlinear remainder is at most $\varepsilon$, and $t^* = (\varepsilon + \gamma + \sqrt{\gamma^2 + 3.84 + 2\|r_0\|\varepsilon + 2\gamma\varepsilon})/\sqrt{I_j}$ is smaller than the distances $d_j^-$ and $d_j^+$ from the optimum to the bounds of $\log|\theta_j|$, then $\theta_j$ is practically identifiable. With $\gamma = \varepsilon = 0$, $t^*$ reduces to the Wald half-width $\sqrt{3.84/I_j}$, and the linear version of the condition, used in Results, is $\sqrt{3.84/I_j} < \min(d_j^-, d_j^+)$; the remainder can be evaluated only at points, so this condition was checked at the computed profile points. *Hypotheses over the admissible set.* Because the theorem requires the remainder bound over the whole admissible range, not only at profile points, we proved two variants (`practically_identifiable_grad_tail` and `practically_identifiable_rel_tail`) that need the bound only on the tail of the profile beyond the guaranteed half-width, either in absolute form ($\le \varepsilon$) or relative to the linear change ($\le \eta\|A\delta\|$ with $\eta < 1$). For every selected parameter and side, the tail bounds were estimated over the region of the free parameters within a factor $r \in \{10, 5, 3, 2, 1.5\}$ of the optimum (and within the bounds of the study) from 512 Sobol points, the profile points in the region and an adversarial search (L-BFGS-B maximisation of the remainder, or of the ratio, on the tail, started from the worst sampled points and repeated until the half-width was stable); a parameter satisfies the hypotheses in the region when one of the two conditions holds on both sides. As a check of the profile resolution, all profiles were recomputed with 20 instead of 10 displacements per side, and the profiles of the subsets selected on the training outputs were computed with the same procedure. Because only the selected and calibration parameters are free, the profiles of a problem took a median of 34 s. *Subsets of other criteria.* The subsets chosen by Brun's collinearity index, Morris screening, pivoted QR and profile likelihood (Comparison with other selection criteria) were profiled with exactly the same procedure: same reduced problem (selected and calibration parameters free, the others at $\theta_0$), all measurements, the same ten displacements per side, refits and bounds, and the same classification (`scripts/run_identifiability_other_criteria.sh`). Where a criterion chose our subset, our profiles were used (10, 6, 2 and 1 problems for Brun, Morris, profile likelihood and pivoted QR). The comparison is made on the 27 problems in which every criterion has a subset (Bertozzi excluded, as in the comparison of admissibility); for each criterion we report the identifiable parameters and the problems in which every selected parameter is identifiable and the subset is admissible at $\pm 5\%$, and compare the latter with ours by the exact McNemar test with Holm adjustment over the four criteria (Statistics). These profiles took 41 h of processor time (Brun 4.7 h, Morris 13.1 h, profile likelihood 8.6 h, pivoted QR 14.4 h). The complete mathematical derivation, the profiles of every problem and the values of all

quantities for every parameter are given in Supplementary Note *Identifiability of the selected parameters*. The profiles were computed for all 31 problems with an admissible subset: those of Bachmann, Chen, Isensee and Smith on a 256-core server, that of Raimundez on a 128-core server, and the other 26 on an 8-core workstation (libRoadRunner 2.10.0 on the servers and 2.9.2 on the workstation).

## Robustness patterns

For each admissible problem, the final subset was kept fixed and $e_{\mathrm{fit}}$ was evaluated at perturbation sizes $\varepsilon \in \{0.01, 0.05, 0.10, 0.20, 0.30, 0.40, 0.50\}$ with 15 scenarios per level. The fifteen directions $u^{(k)} \in [-1, 1]^p$, with independent uniform components, are drawn once per problem with a fixed seed (42) and reused at every level, $\theta^{(k)} = \theta_0 \odot (1 + \varepsilon u^{(k)})$; the directions are stored with the results. To check that the patterns do not depend on these particular directions, the evaluation was repeated with new directions at each size (seed 42 at $\pm 5\%$, which reproduces the selection, and $42 + 1000i$ at the $i$-th other level). Let $\bar{e}(\varepsilon)$ be the median $e_{\mathrm{fit}}$ at level $\varepsilon$ and $\varepsilon_f$ the smallest level with $\bar{e}(\varepsilon_f) > \tau_e$. The pattern is A if no such level exists; C (irregular) if $\bar{e}(\varepsilon) \leq \tau_e$ again at a level above $\varepsilon_f$; otherwise B if $\varepsilon_f \geq 0.20$ and D if $\varepsilon_f \leq 0.10$; the pass rate is the fraction of valid scenarios with $e_{\mathrm{fit}} \leq \tau_e$. As in the selection, $e_{\mathrm{fit}}$ was computed to first order for the problems marked in Table 1. For Chen, between 8 and 13 of the 15 scenarios per level could be simulated in the run shown in Fig. 5 (7 to 13 over all our runs, because some simulations fail in a run-dependent way); the remaining scenarios failed to integrate and were discarded. A scenario is valid when it could be simulated and its output change exceeds the numerical noise floor; all fifteen scenarios were valid in 28 of the 31 problems, and only 8, 4 and 2 in Chen, Isensee and Perelson. Among the 26 problems of pattern A, 9 pass every scenario at every level, whereas 7 (Borghans, Bruno, Fujita, Liu, Sneyd, Zhao, Zheng) fall to a pass rate of 60–73% at some level. The full distribution of the error for every problem and level is given in Supplementary Tables S11 and S12 and Supplementary Fig. S4.

## Computational cost

The cost of the selection is dominated by the computation of $J$ ($2p$ simulations) and of the stage-2 ordering ($p$ simulations); the selection, the conditioning tests and the first-order $e_{\mathrm{fit}}$ only use linear algebra on $J$ (Supplementary Note *Computational cost*). The computation time of the complete method was measured by running it from scratch, without cached results, with the sensitivity computations distributed over 16–64 cores of a 128-core server shared with other computations, so the times may be inflated by resource contention. Times are given as elapsed (wall-clock) time when the computation of a problem ran in parallel, and as processor time, the sum over tasks, when it consisted of independent single-core tasks (scenarios, seeds, profile points). For Armistead, Lang and Raimundez the times were measured with a slightly different scenario set; the sensitivity simulations, which dominate the cost, are the same. Chen, with 35 s per simulation, was the slowest problem (33 min). In the six problems selected with the first-order $e_{\mathrm{fit}}$, the nonlinear re-estimation of the final subset in the 15 scenarios took 4 min (Lucarelli) to 12.7 h (Chen) of processor

time (Alkan 8 min, Bachmann 23 min, Isensee 3.9 h, Raimundez 4.1 h). Profile likelihood was much slower on problems with expensive simulations (Alkan 27 min against 25 s for our method; Zheng 5 min against 14 s; Isensee 5.2 h, or 110 h of processor time, against 10 min). For Bachmann (90 free parameters in the profile, 84 in the synthetic experiment after removing those that act only in test conditions; Supplementary Table S3) and Lucarelli (72), profile likelihood took 70 and 11 min with the profile points distributed over 60 cores (40 and 9 h of computing time), against 10 and 35 s for our method; Chen and Raimundez were processed by our method in 33 and 2 min. Selecting a subset took less than ten minutes in 32 of the 34 problems. All analyses reported in this paper together took about 697 h of processor time, the sum over all cores of the times recorded by each task in the final runs (only the final runs are counted, and the compilation of the Lean proofs is not included); because the analyses ran in parallel on servers with 88–256 cores and on an 8-core workstation, the elapsed time was much shorter. 87% of the processor time was spent in Isensee, Chen, Raimundez, Smith and Bachmann; the largest items were profile likelihood of the full model (193 h), the identifiability profiles (172 h: 26 h for our subsets, 102 h for the check with 20 points per side, 3 h for the subsets selected on the training outputs and 41 h for the subsets of the other criteria), the tests of the subsets of all criteria (86 h, and 17 h more with the calibration parameters fixed) and the synthetic experiments (130 h over all amplitudes, families and criteria). The check of the identifiability hypotheses over the admissible set took a further 3 h of elapsed time on an 88-core server (Supplementary Note *Computing time of the analyses*). On state trajectories, the 4,231-parameter Froehlich model required 12,693 simulations and 2.6 h for the sensitivities; for models of this size on measured outputs, forward sensitivity equations[49] compute $J$ in a single integration of an augmented system, whereas adjoint methods are efficient for gradients of a scalar objective but not for the full output Jacobian needed here.

## Formal verification

The formal development uses Lean 4 (version 4.23.0) and Mathlib[37, 38] and has two parts (Table 3 and Supplementary Fig. S3). The first proves, for any model that satisfies stated hypotheses, the mathematical results on which each step of the method relies (Supplementary Table S4, steps 1–22). The second establishes those hypotheses for concrete benchmark models. For this, each PEtab problem was translated automatically into a Lean file (`lean_translation/sbml_to_lean.py` of the reproducibility package): the rate laws and stoichiometry of the SBML model and the nominal parameter values and initial conditions of the PEtab tables are written as Lean expressions in which every number is an exact rational. Verified checkers, Boolean functions proved sound in Lean, then test that every rate law is defined and smooth on a domain containing the trajectory, that concentrations cannot become negative (quasi-positivity) and that the growth of the vector field is bounded so that the solution exists on the simulated horizon (linear growth, or a Riccati bound on a finite horizon for Crauste). The checkers are evaluated by the Lean kernel itself (`decide +kernel`), so their results are part of the proof, and each model ends with a theorem stating that it is smooth, that its solution exists and

remains non-negative on the horizon, and that the solution is differentiable with respect to the parameters, which are the hypotheses of steps 14–22. This was achieved for 28 benchmark models (27 of the 34 problems analysed, and Froehlich); the reasons why the other seven problems were not verified are given in Supplementary Note *Verified benchmark models*. Because the theorems are only as relevant as the translation is faithful, the translated vector fields, parameter values and initial values were compared independently with the SBML models evaluated by libRoadRunner at the nominal state and at five random states; they agree to rounding error (largest relative difference $10^{-12}$), and changing a single number of a Lean file by 1% is detected. All files compile from scratch without `sorry`; `#print axioms` reports only the standard axioms of Lean's logic (`propext`, `Classical.choice`, `Quot.sound`) for all 309 theorems and lemmas and the 28 model theorems, and every module is accepted by the independent checker `lean4checker`. The file `NonVacuity.lean` gives, for each central theorem, a concrete instance in which every hypothesis is proved (including a nonlinear residual for the identifiability theorems), so that no central statement is vacuous, and `DEFINITIONS_AUDIT.md` compares each Lean definition with the paper and the code. The numerical analyses of this paper ($J$, $R_{\mathrm{var}}$, $\kappa$, VIF, $\cos\Delta$, $e_{\mathrm{fit}}$ and the validation) were carried out in Python in double precision; a formal conclusion applies to a numerical result when the model's hypotheses are verified in Lean and the computed quantities satisfy the inequalities that the corresponding theorem requires (Supplementary Table S5).

*Further verified statements.* Besides Theorems 1 and 2, the first part proves that, to first order, fixing the non-selected parameters changes the outputs by at most $\sqrt{(1 - R_{\mathrm{var}}(S))\, E}\, \|\Delta\theta\|$, where $E = \sum_j E_j$ (by the Cauchy–Schwarz inequality applied to the discarded columns of $J$), which is the error bound behind level 3; that re-estimating more parameters never increases the fitting error and that $e_{\mathrm{fit}} \le e_{\mathrm{rel}}$, so that the admissibility criterion is monotone when the re-estimation reaches its optimum (with a local optimiser this holds only approximately, which is why monotonicity was also checked empirically); that the greedy selection keeps $\kappa$ and VIF below their thresholds and that discarding a rejected candidate permanently is correct; that bounds on $\kappa$ and VIF of the unit-length columns imply a unique and stable solution of the linearised least-squares problem for the retained parameters (which is not a proof of identifiability of the nonlinear model); that error bounds on the states transfer to the measured outputs, including logarithmic transformations; and that the first-order $e_{\mathrm{fit}}$ differs from the exact one by a term of second order in the absolute error (which does not by itself bound the relative error used in the decision). Explicit versions of the error bound and of the cosine for finite perturbations are also proved; they bound the second-order remainder through a Lipschitz constant $L_c$ computed from bounds on the first and second derivatives of the vector field. Evaluated for the verified models, $L_c$ ranged from about $10^2$ to $10^{24}$, too large to certify any scenario at $\pm 5\%$, so these finite certificates are of theoretical value only. Theorem `practically_identifiable_grad` (step 32) is stated in the Methods section *Identifiability of the selected parameters*; its hypothesis on the nonlinear remainder is required for all parameter values in the admissible range, whereas we evaluated it at the computed profile points. In seven problems (Armistead, Blasi, Boehm, Bruno,

Rahman, Sneyd, Zheng), every selected parameter satisfied these hypotheses at all evaluated points and the model itself is verified in Lean; in the other 24 admissible problems identifiability rests on the computed profiles alone. Supplementary Table S43 lists, for every claim of the Results that refers to the formal development, the Lean theorem, its hypotheses and the numerical quantities that instantiate them.

## Use of large language models

An AI assistant (Claude, Anthropic) was used to help write and debug analysis code and Lean proofs and to edit the manuscript text. All code, proofs and results were checked by the authors, who take full responsibility for the content.

**Data availability.** The PEtab benchmark problems analysed in this study are publicly available in the Benchmark-Models-PEtab collection (https://github.com/Benchmarking-Initiative/Benchmark-Models-PEtab, commit `fcbddf1`). All results generated in this study (selected subsets, external validation, synthetic prediction experiments, comparison with other selection criteria including the profile likelihoods, identifiability of the selected parameters, robustness, single-parameter subsets, sensitivity to the method settings, simulator check, computation times and the checks of the Lean models), one file per benchmark problem with a data dictionary, are included in the reproducibility package archived on Zenodo (doi:10.5281/zenodo.23338175) and in the folder `reproducibility/` of https://github.com/monicammr/archivos.

**Code availability.** The reproducibility package (Zenodo, doi:10.5281/zenodo.23338175; folder `reproducibility/` of https://github.com/monicammr/archivos) contains the Python code of the method and of every analysis, the programs that produce every figure and table of the paper from the stored results, the translator of the SBML models into Lean, and all Lean 4 proofs (folder `lean_proofs/`, a self-contained project that compiles with a single command, `./verify.sh`, and checks that no proof uses `sorry` or additional axioms). Its manual (`README.md`) describes three independent levels of checking: (A) regenerating every figure, table and number of the paper from the stored results, without simulation, in a few minutes, which shows that the paper follows from those results; (B) compiling all proofs from scratch with `./verify.sh`, in 20 min to 2 h, which shows that every theorem is proved with no unproved assumption; and (C) re-running the analyses from the public PEtab problems and comparing the new results with the stored ones, which shows that the stored results are those the code produces (minutes per problem for the method itself, about 697 h of processor time for all analyses of the paper). Supplementary Note *Reproducibility package* describes each level in detail.

**Acknowledgements.** The authors thank Escuela Superior Politécnica del Litoral (ESPOL) for institutional support.

**Author contributions.** M.M.R. and C.S. conceived the study and designed the method. M.M.R. wrote the software and the formal proofs, performed the analyses and wrote the first draft. C.S., C.M. and A.A. supervised the work, contributed to

the validation of the results and revised the manuscript. All authors reviewed and approved the final manuscript.

**Funding.** This research received no external funding.

**Ethics declarations. Ethics approval and consent to participate:** Ethics approval was not required for this secondary analysis of publicly available, de-identified data. No new data from human participants were collected for this study. All analyses used publicly available, previously published benchmark models and data (PEtab benchmark collection); some of these data are human-derived, such as the published viral-load measurements of a patient (Perelson problem) and aggregated epidemiological case counts, and were analysed secondarily, in the de-identified form in which they were published. No identifiable participant information was accessed. Ethical approval and informed consent for the original studies were the responsibility of the corresponding original investigations.
**Consent for publication:** Not applicable.

**Additional information. Competing interests:** The authors declare no competing interests.
**Correspondence** and requests for materials should be addressed to M.M.R.
**Supplementary information** accompanies this paper (Supplementary Notes, Tables and Figures); Supplementary Note *Concluding remarks* summarises its contents.

# References

[1] Kholodenko, B. N. Cell-signalling dynamics in time and space. *Nature Reviews Molecular Cell Biology* **7**, 165–176 (2006).

[2] Bellman, R. & Åström, K. J. On structural identifiability. *Mathematical Biosciences* **7**, 329–339 (1970).

[3] Chiş, O.-T., Banga, J. R. & Balsa-Canto, E. Structural identifiability of systems biology models: A critical comparison of methods. *PLOS ONE* **6**, e27755 (2011).

[4] Villaverde, A. F., Barreiro, A. & Papachristodoulou, A. Structural identifiability of dynamic systems biology models. *PLOS Computational Biology* **12**, e1005153 (2016).

[5] Miao, H., Xia, X., Perelson, A. S. & Wu, H. On identifiability of nonlinear ODE models and applications in viral dynamics. *SIAM Rev.* **53**, 3–39 (2011).

[6] Wieland, F.-G., Hauber, A. L., Rosenblatt, M., Tönsing, C. & Timmer, J. On structural and practical identifiability. *Curr. Opin. Syst. Biol.* **25**, 60–69 (2021).

[7] Raue, A. *et al.* Structural and practical identifiability analysis of partially observed dynamical models by exploiting the profile likelihood. *Bioinformatics* **25**, 1923–1929 (2009).

[8] Raue, A., Kreutz, C., Theis, F. J. & Timmer, J. Joining forces of Bayesian and frequentist methodology: A study for inference in the presence of non-identifiability. *Philosophical Transactions of the Royal Society A* **371**, 20110544 (2013).

[9] Kreutz, C., Raue, A. & Timmer, J. Likelihood based observability analysis and confidence intervals for predictions of dynamic models. *BMC Systems Biology* **6**, 120 (2012).

[10] Hass, H. *et al.* Benchmark problems for dynamic modeling of intracellular processes. *Bioinformatics* **35**, 3073–3082 (2019).

[11] Gutenkunst, R. N. *et al.* Universally sloppy parameter sensitivities in systems biology models. *PLOS Computational Biology* **3**, e189 (2007).

[12] Transtrum, M. K. *et al.* Perspective: Sloppiness and emergent theories in physics, biology, and beyond. *J. Chem. Phys.* **143**, 010901 (2015).

[13] Monsalve-Bravo, G. M. *et al.* Analysis of sloppiness in model simulations: Unveiling parameter uncertainty when mathematical models are fitted to data. *Sci. Adv.* **8**, eabm5952 (2022).

[14] Saltelli, A., Tarantola, S., Campolongo, F. & Ratto, M. *Sensitivity analysis in practice: A guide to assessing scientific models* (Wiley, 2004).

[15] Morris, M. D. Factorial sampling plans for preliminary computational experiments. *Technometrics* **33**, 161–174 (1991).

[16] Maiwald, T. *et al.* Driving the model to its limit: Profile likelihood based model reduction. *PLOS ONE* **11**, e0162366 (2016).

[17] Eisenberg, M. C. & Hayashi, M. A. L. Determining identifiable parameter combinations using subset profiling. *Math. Biosci.* **256**, 116–126 (2014).

[18] Belsley, D. A., Kuh, E. & Welsch, R. E. *Regression diagnostics: Identifying influential data and sources of collinearity* (Wiley, 1980).

[19] Weijers, S. R. & Vanrolleghem, P. A. A procedure for selecting best identifiable parameters in calibrating activated sludge model no. 1 to full-scale plant data. *Water Sci. Technol.* **36**, 69–79 (1997).

[20] Cintrón-Arias, A., Banks, H. T., Capaldi, A. & Lloyd, A. L. A sensitivity matrix based methodology for inverse problem formulation. *J. Inverse Ill-Posed Probl.* **17**, 545–564 (2009).

[21] Brun, R., Reichert, P. & Künsch, H. R. Practical identifiability analysis of large environmental simulation models. *Water Resources Research* **37**, 1015–1030 (2001).

[22] Gábor, A., Villaverde, A. F. & Banga, J. R. Parameter identifiability analysis and visualization in large-scale kinetic models of biosystems. *BMC Systems Biology* **11**, 54 (2017).

[23] Yao, K. Z., Shaw, B. M., Kou, B., McAuley, K. B. & Bacon, D. W. Modeling ethylene/butene copolymerization with multi-site catalysts: Parameter estimability and experimental design. *Polym. React. Eng.* **11**, 563–588 (2003).

[24] Li, R., Henson, M. A. & Kurtz, M. J. Selection of model parameters for off-line parameter estimation. *IEEE Trans. Control Syst. Technol.* **12**, 402–412 (2004).

[25] Lund, B. F. & Foss, B. A. Parameter ranking by orthogonalization—Applied to nonlinear mechanistic models. *Automatica* **44**, 278–281 (2008).

[26] Quaiser, T. & Mönnigmann, M. Systematic identifiability testing for unambiguous mechanistic modeling – application to JAK-STAT, MAP kinase, and NF-$\kappa$B signaling pathway models. *BMC Syst. Biol.* **3**, 50 (2009).

[27] Golub, G. H., Klema, V. & Stewart, G. W. Rank degeneracy and least squares problems. *NBER Working Paper* **165**; https://doi.org/10.3386/w0165 (1977).

[28] Vélez-Reyes, M. & Verghese, G. C. Subset selection in identification, and application to speed and parameter estimation for induction machines. In *Proceedings of the 4th IEEE Conference on Control Applications*, 991–997 (1995).

[29] Golub, G. H. & Van Loan, C. F. *Matrix Computations* (4th edn, Johns Hopkins University Press, 2013).

[30] Olufsen, M. S. & Ottesen, J. T. A practical approach to parameter estimation applied to model predicting heart rate regulation. *J. Math. Biol.* **67**, 39–68 (2013).

[31] Bandura, A. *Social Foundations of Thought and Action: A Social Cognitive Theory* (Prentice-Hall, 1986).

[32] Martín, C. A. *et al.* A dynamical systems model of social cognitive theory. In *2014 American Control Conference (ACC)*, 2407–2412 (IEEE, 2014).

[33] Riley, W. T. *et al.* Development of a dynamic computational model of social cognitive theory. *Transl. Behav. Med.* **6**, 483–495 (2016).

[34] Martín, C. A. *et al.* Development of a control-oriented model of social cognitive theory for optimized mHealth behavioral interventions. *IEEE Trans. Control Syst. Technol.* **28**, 331–346; https://doi.org/10.1109/TCST.2018.2873538 (2020).

[35] Miranda R., M., Salazar L., C., Martín M., C. & Aguirre A., A. Hybrid estimation of dynamic behavioral models using differential evolution and grey-box refinement. In *2025 IEEE Ninth Ecuador Technical Chapters Meeting (ETCM)*, 1–7; https://doi.org/10.1109/ETCM67548.2025.11304524 (IEEE, 2025).

[36] Schmiester, L. *et al.* PEtab—interoperable specification of parameter estimation problems in systems biology. *PLOS Computational Biology* **17**, e1008646 (2021).

[37] de Moura, L. & Ullrich, S. The Lean 4 theorem prover and programming language. In *Automated Deduction – CADE 28*, Lecture Notes in Computer Science **12699**, 625–635 (Springer, 2021).

[38] Mathlib Community. The Lean mathematical library. In *Proceedings of the 9th ACM SIGPLAN International Conference on Certified Programs and Proofs (CPP 2020)*, 367–381; https://doi.org/10.1145/3372885.3373824 (ACM, 2020).

[39] Liepe, J. *et al.* A framework for parameter estimation and model selection from experimental data in systems biology using approximate Bayesian computation. *Nature Protocols* **9**, 439–456 (2014).

[40] Transtrum, M. K. & Qiu, P. Model reduction by manifold boundaries. *Physical Review Letters* **113**, 098701 (2014).

[41] Moore, B. C. Principal component analysis in linear systems: Controllability, observability, and model reduction. *IEEE Transactions on Automatic Control* **26**, 17–32 (1981).

[42] Antoulas, A. C. *Approximation of large-scale dynamical systems* (SIAM, 2005).

[43] Loos, C., Krause, S. & Hasenauer, J. Hierarchical optimization for the efficient parametrization of ODE models. *Bioinformatics* **34**, 4266–4273 (2018).

[44] Fröhlich, F., Kaltenbacher, B., Theis, F. J. & Hasenauer, J. Scalable parameter estimation for genome-scale biochemical reaction networks. *PLOS Computational Biology* **13**, e1005331 (2017).

[45] Transtrum, M. K. & Qiu, P. Optimal experiment selection for parameter estimation in biological differential equation models. *BMC Bioinformatics* **13**, 181 (2012).

[46] Welsh, C. *et al.* libRoadRunner 2.0: a high performance SBML simulation and analysis library. *Bioinformatics* **39**, btac770 (2023).

[47] Virtanen, P. *et al.* SciPy 1.0: fundamental algorithms for scientific computing in Python. *Nat. Methods* **17**, 261–272 (2020).

[48] Villaverde, A. F., Fröhlich, F., Weindl, D., Hasenauer, J. & Banga, J. R. Benchmarking optimization methods for parameter estimation in large kinetic models. *Bioinformatics* **35**, 830–838 (2019).

[49] Fröhlich, F. *et al.* AMICI: High-performance sensitivity analysis for large ordinary differential equation models. *Bioinformatics* **37**, 3676–3677 (2021).